



CS 4495 Computer Vision

Finding 2D Shapes and the Hough Transform

Aaron Bobick

School of Interactive
Computing



Administrivia

- Today: Modeling Lines and Finding them
- Problem set 1 is posted.
 - You **can** use Matlab edge operators
 - You **cannot** use Matlab Hough methods.
 - Due Sunday, Sept 8th 11:55pm.

Now some “real” vision...

- So far, we applied operators/masks/kernels to images to produce new image




Image processing: $F : I(x, y) \longrightarrow I'(x, y)$

- Now real vision:

$$F : I(x, y) \longrightarrow \text{good stuff}$$

Fitting a model

- Want to associate a model with observed features



[Fig from Marszalek & Schmid, 2007]

For example, the model could be a line, a circle, or an arbitrary shape.


Fitting

- Choose a ***parametric model*** to represent a set of features
- Membership criterion is not local
 - Can't tell whether a point in the image belongs to a given model just by looking at that point
- Three main questions:
 1. What model represents this set of features best?
 2. Which of several model instances gets which feature?
 3. How many model instances are there?
- Computational complexity is important
 - It is infeasible to examine every possible set of parameters and every possible combination of features

Example: Line fitting



- Why fit lines?

Many objects characterized by presence of straight lines



- Wait, why aren't we done just by running edge detection?

Difficulty of line fitting





- **Extra** edge points (clutter), multiple models:
 - which points go with which line, if any?
- Only some parts of each line detected, and some parts are **missing**:
 - how to find a line that bridges missing evidence?
- **Noise** in measured edge points, orientations:
 - how to detect true underlying parameters?

Voting


- It's not feasible to check all combinations of features by fitting a model to each possible subset.
- **Voting** is a general technique where we let the features vote for all models that are compatible with it.
 - Cycle through features, cast votes for model parameters.
 - Look for model parameters that receive a lot of votes.
- Noise & clutter features will cast votes too, but typically their votes should be inconsistent with the majority of “good” features.
- Ok if some features not observed, as model can span multiple fragments.

Fitting lines

- Given points that belong to a line, what is the line?
- How many lines are there?
- Which points belong to which lines?
- **Hough Transform** is a voting technique that can be used to answer all of these
 - Main idea:
 - 1. Record all possible lines on which each edge point lies.
 - 2. Look for lines that get many votes.




Finding lines in an image: Hough space



Connection between image (x,y) and Hough (m,b) spaces

- A line in the image corresponds to a point in Hough space
- To go from image space to Hough space:
 - given a set of points (x,y) , find all (m,b) such that $y = mx + b$


Finding lines in an image: Hough space



Connection between image (x,y) and Hough (m,b) spaces

- A line in the image corresponds to a point in Hough space
- To go from image space to Hough space:
 - given a set of points (x,y) , find all (m,b) such that $y = mx + b$
- What does a point (x_0, y_0) in the image space map to?
 - Answer: the solutions of $b = -x_0m + y_0$
 - this is a line in Hough space

Finding lines in an image: Hough *transform*



What are the line parameters for the line that contains both (x_0, y_0) and (x_1, y_1) ?

- It is the intersection of the lines $b = -x_0 m + y_0$ and $b = -x_1 m + y_1$

Finding lines: Hough algorithm





image space




Hough (parameter) space

- How can we use this to find the most likely parameters (m,b) for the most prominent line in the image space?
- Let each edge point in image space vote for a set of possible parameters in Hough space
- Accumulate votes in discrete set of bins; parameters with the most votes indicate line in image space.

Polar representation for lines

Issues with usual (m,b) parameter space: can take on infinite values, undefined for vertical lines.



d : perpendicular distance
from line to origin

θ : angle the perpendicular
makes with the x-axis

$$x \cos \theta - y \sin \theta = d$$

Point in image space \rightarrow sinusoid segment in Hough space

Hough transform algorithm


Using the polar parameterization:

$$x \cos \theta - y \sin \theta = d$$

Basic Hough transform algorithm

1. Initialize $H[d, \theta] = 0$
2. for each edge point $I[x, y]$ in the image
for $\theta = 0$ to 180 // some quantization; not 2π ?
 $d = x \cos \theta - y \sin \theta$ // maybe negative
 $H[d, \theta] += 1$
3. Find the value(s) of (d, θ) where $H[d, \theta]$ is maximum
4. The detected line in the image is given by $d = x \cos \theta - y \sin \theta$


H: accumulator array (votes)




Space complexity? k^n (n dimensions, k bins each)

Time complexity (in terms of number of voting elements)?

Example: Hough transform for straight lines



**Image space
edge coordinates**




Votes



Bright value = high vote count
Black = no votes

Example: Hough transform for straight lines




Square :





Example: Hough transform for straight lines




Hough demo..






Showing longest segments
found

Impact of noise on Hough




**Image space
edge coordinates**




Votes

What difficulty does this present for an implementation?

Impact of noise on Hough



**Image space
edge coordinates**



Votes

Here, everything appears to be “noise”, or random edge points, but we still see peaks in the vote space.

Extensions

- **Extension 1:** Use the image gradient

- same
- for each edge point $I[x,y]$ in the image
 - θ = gradient at (x,y)

$$d = x \cos \theta - y \sin \theta$$

- $H[d, \theta] += 1$
- same
- same
- (Reduces degrees of freedom)

- Extension 2


- give more votes for stronger edges

- Extension 3

- change the sampling of (d, θ) to give more/less resolution

- Extension 4

- The same procedure can be used with circles, squares, or any other shape



$$\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

Hough transform for circles

- Circle: center (a,b) and radius r

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- For a fixed radius r , unknown gradient direction





Image space




Hough space a

Hough transform for circles

- Circle: center (a,b) and radius r

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- For a fixed radius r , unknown gradient direction




Hough transform for circles

- Circle: center (a,b) and radius r

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- For an unknown radius r , unknown gradient direction



Hough transform for circles

- Circle: center (a,b) and radius r

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- For an unknown radius r , unknown gradient direction





Image space



Hough space

Hough transform for circles

- Circle: center (a,b) and radius r

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- For an unknown radius r , **known** gradient direction





Image space



Hough space

Hough transform for circles

For every edge pixel (x,y) :

For each possible radius value r .

For each possible gradient direction θ :
%% or use estimated gradient

$$a = x - r \cos(\theta)$$

$$b = y + r \sin(\theta)$$

$$H[a,b,r] += 1$$

end

end


Example: detecting circles with Hough




Crosshair indicates results of Hough transform,
bounding box found via motion differencing.

Example: detecting circles with Hough


Original



Edges




Votes: Penny




Note: a different Hough transform (with separate accumulators) was used for each circle radius (quarters vs. penny).

Example: detecting circles with Hough


Original



Edges



Votes: Quarter



Voting: practical tips

- Minimize irrelevant tokens first (take edge points with significant gradient magnitude)
- Choose a good grid / discretization
 - Too coarse: large votes obtained when too many different lines correspond to a single bucket
 - Too fine: miss lines because some points that are not exactly collinear cast votes for different buckets
- Vote for neighbors, also (smoothing in accumulator array)
- Utilize direction of edge to reduce free parameters by 1
- To read back which points voted for “winning” peaks, keep tags on the votes.

Hough transform: pros and cons

- Pros


- All points are processed independently, so can cope with occlusion
- Some robustness to noise: noise points unlikely to contribute consistently to any single bin
- Can detect multiple instances of a model in a single pass

- Cons

- Complexity of search time increases exponentially with the number of model parameters
- Non-target shapes can produce spurious peaks in parameter space
- Quantization: hard to pick a good grid size

Generalized Hough transform

- What if want to detect arbitrary shapes defined by boundary points and a reference point?



At each boundary point, compute displacement vector: $\mathbf{r} = \mathbf{a} - \mathbf{p}_i$.

For a given model shape: store these vectors in a table indexed by gradient orientation θ .


[Dana H. Ballard, Generalizing the Hough Transform to Detect Arbitrary Shapes, 1980]

Generalized Hough transform

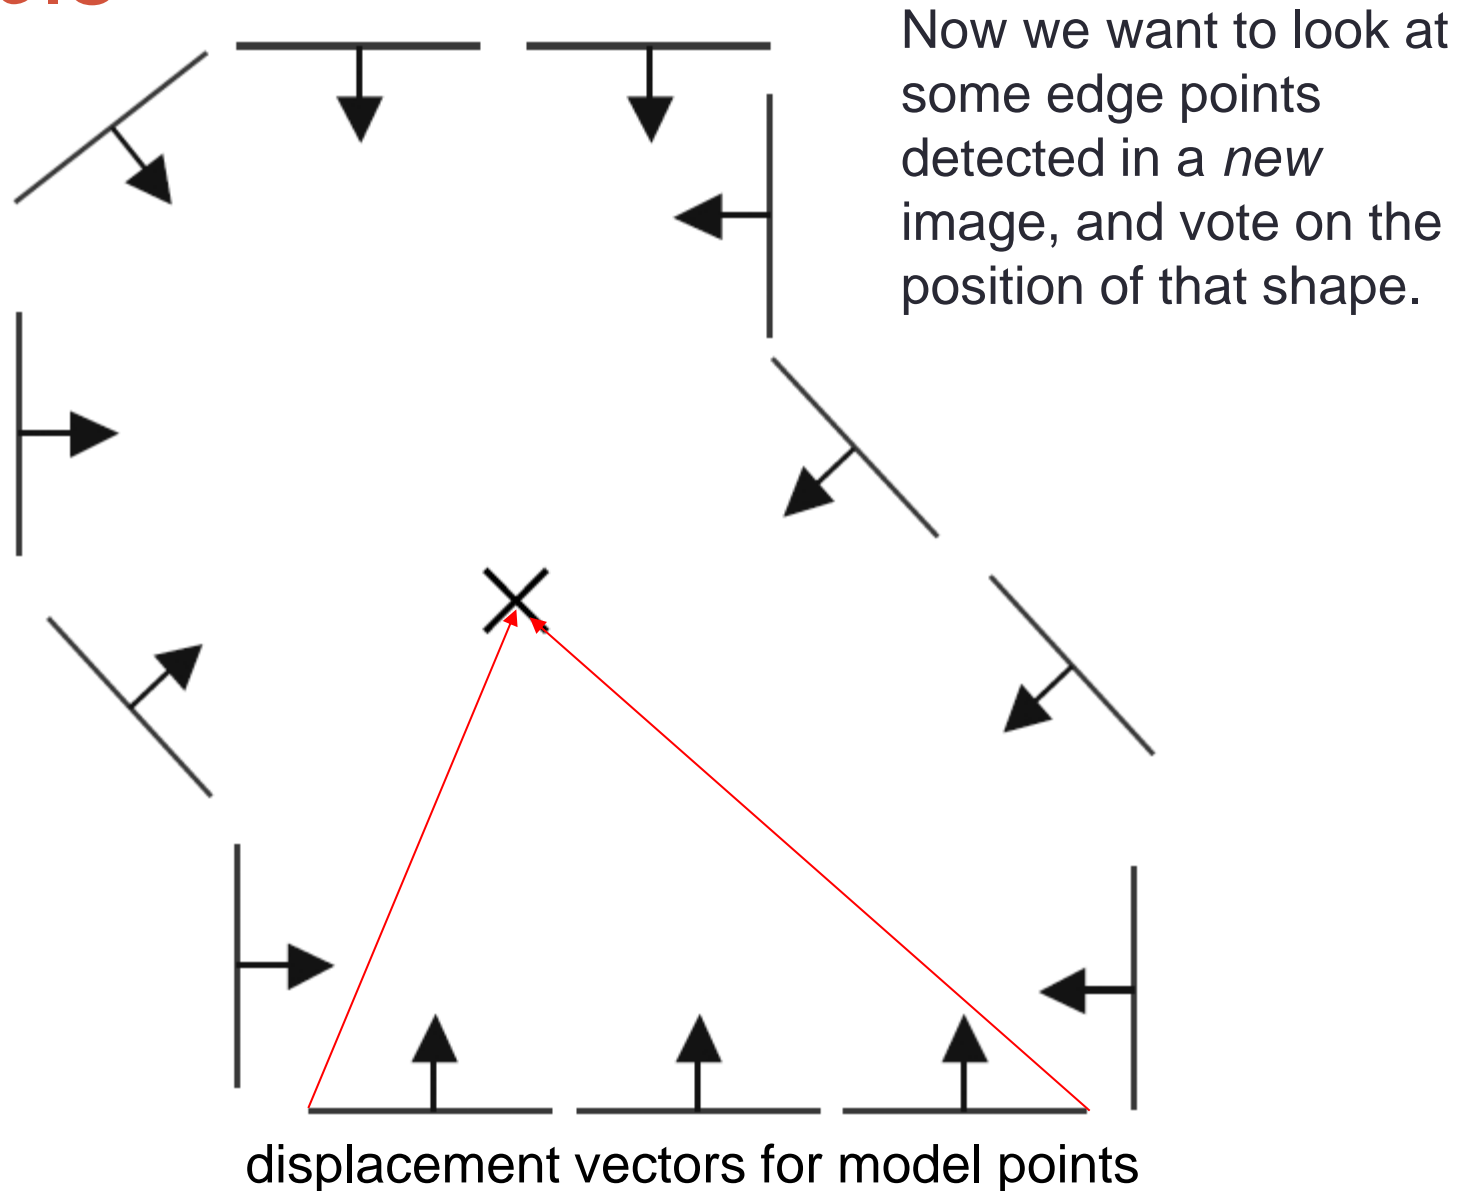
- To detect the model shape in a new image:
- For each edge point
 - For each possible θ^* master orientation
 - Index into table with its gradient orientation $\theta - \theta^*$
 - Use retrieved r vectors to vote for position of reference point
- Peak in this Hough space is reference point with most supporting edges

(In next slides we'll assume only translation is unknown, i.e., orientation and scale are fixed)


Example



Example




Example




range of voting locations for test point


Example




Example




Example



Example




Example




Application in recognition

- Instead of indexing displacements by gradient orientation, index by “visual codeword”



training image



visual codeword with
displacement vectors

B. Leibe, A. Leonardis, and B. Schiele, [Combined Object Categorization and Segmentation with an Implicit Shape Model](#), ECCV Workshop on Statistical Learning in Computer Vision 2004

Application in recognition

- Instead of indexing displacements by gradient orientation, index by “visual codeword”



test image

B. Leibe, A. Leonardis, and B. Schiele, [Combined Object Categorization and Segmentation with an Implicit Shape Model](#), ECCV Workshop on Statistical Learning in Computer Vision 2004

Summary

- Fitting problems require finding any supporting evidence for a model, even within clutter and missing features.
 - associate features with an explicit model
- Voting approaches, such as the Hough transform, make it possible to find likely model parameters without searching all combinations of features.
 - Hough transform approach for lines, circles, ..., arbitrary shapes defined by a set of boundary points, recognition from patches.
- It's Labor Day Weekend – enjoy!