

CS 4495 Computer Vision

Segmentation

Aaron Bobick (slides by Tucker Hermans)
School of Interactive Computing

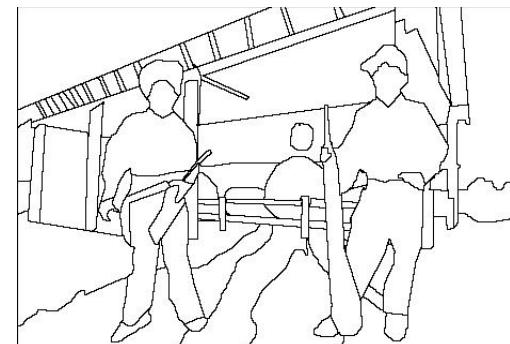
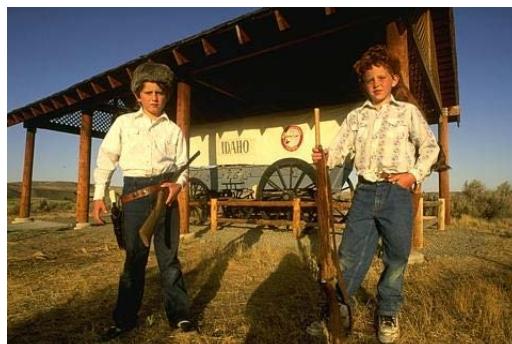
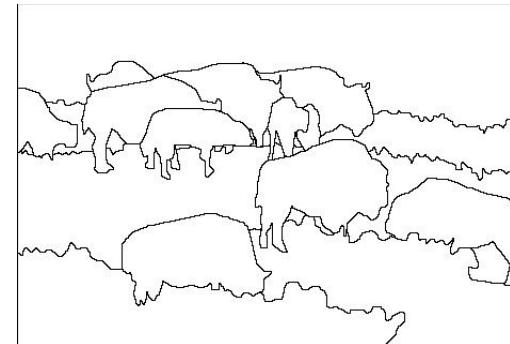


Administrivia

- PS 4: Out but I was a bit late so due date pushed back to Oct 29.
- OpenCV now has real SIFT again (the “notfree” packages). If using Python and OpenCV you should be able to use those calls.
- We’re still investigating SIFT for Python for those *not* using OpenCV.

Why segmentation?

Segmentation of Coherent Regions



Berkeley segmentation database:

<http://www.eecs.berkeley.edu/Research/Projects/CS/vision/grouping/segbench/>

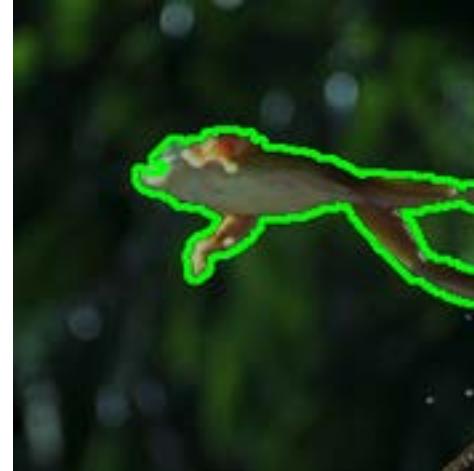
Grouping of Similar Neighbors



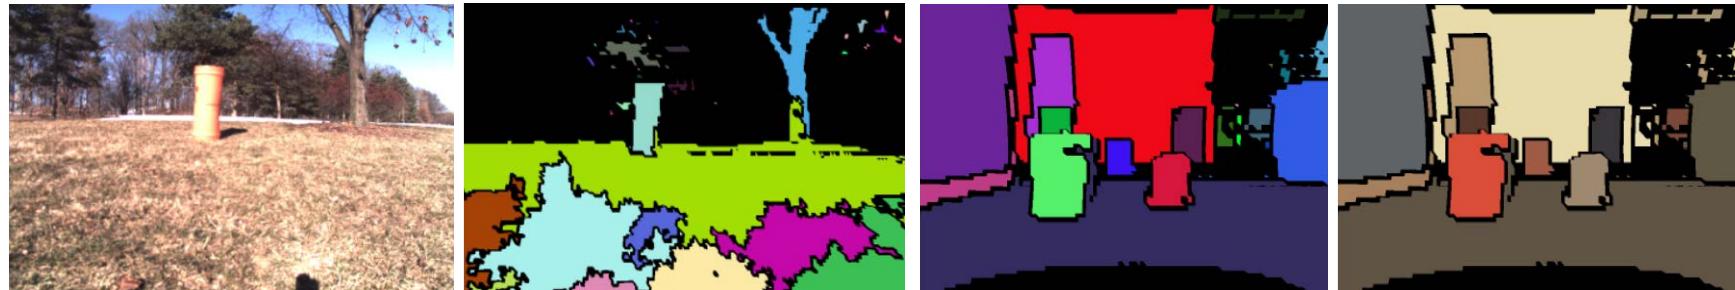
X. Ren and J. Malik. [Learning a classification model for segmentation](#). ICCV 2003.

Figure Ground Segmentation

- Separate the foreground object (figure) from the background (ground)



Extensions Beyond Single Images

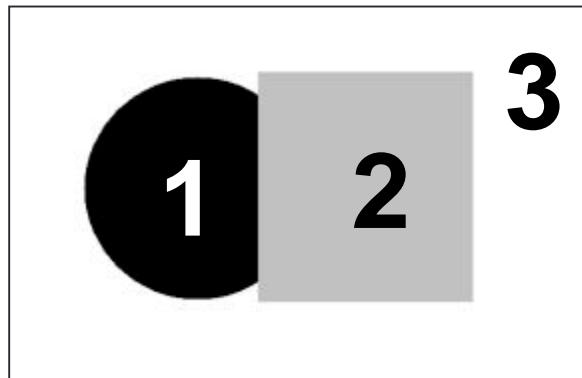


J. Strom, A. Richardson, E. Olson. “Graph-based Segmentation for Colored 3D Laser Point Clouds.” IROS 2010.

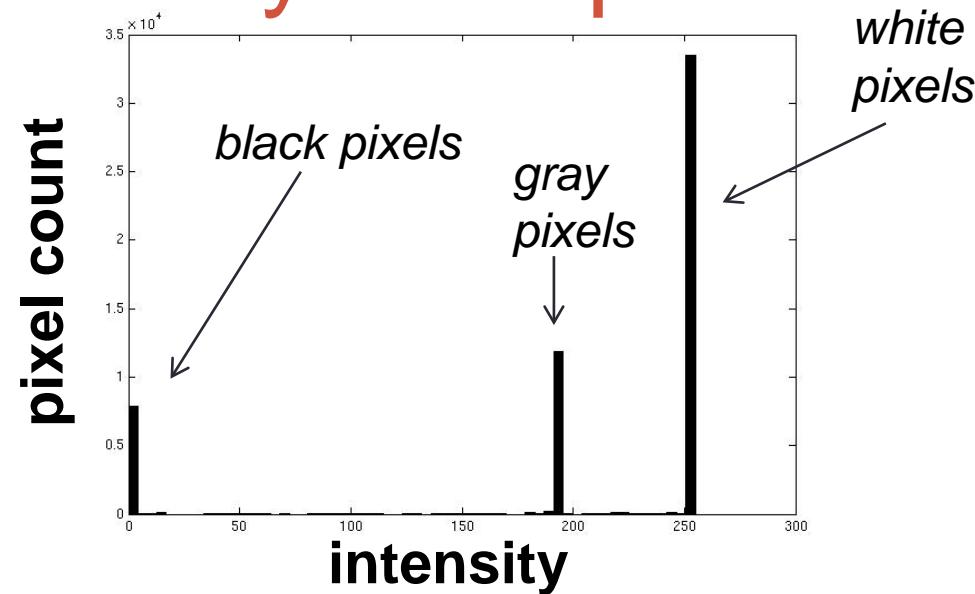


M. Grundmann, V. Kwatra, M. Han, I. Essa. “Efficient Hierarchical Graph-Based Video Segmentation.” CVPR 2010.

Image segmentation: toy example

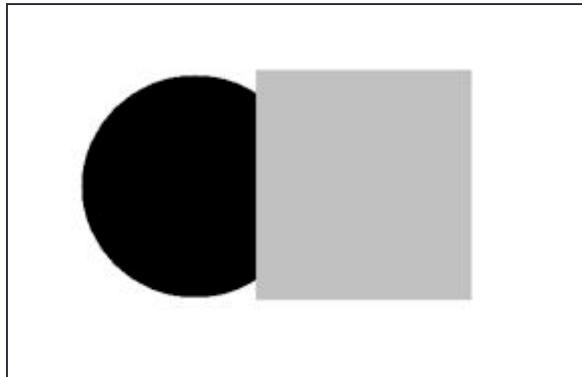


input image

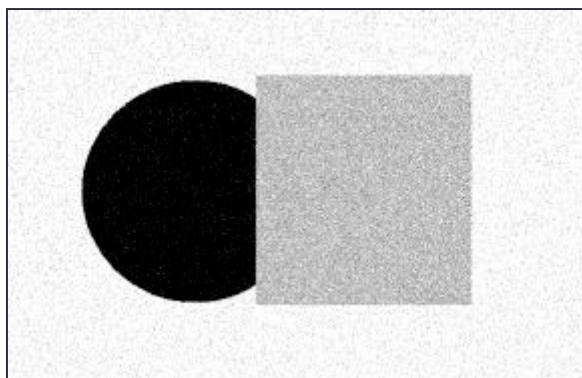
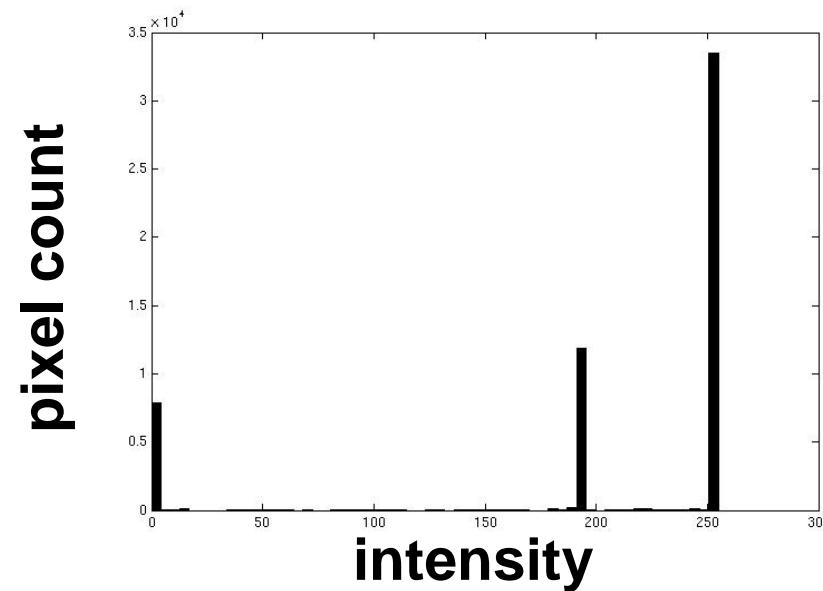


- These intensities define the three groups.
- We could label every pixel in the image according to which of these primary intensities it is.
 - i.e., *segment* the image based on the intensity feature.
- What if the image isn't quite so simple?

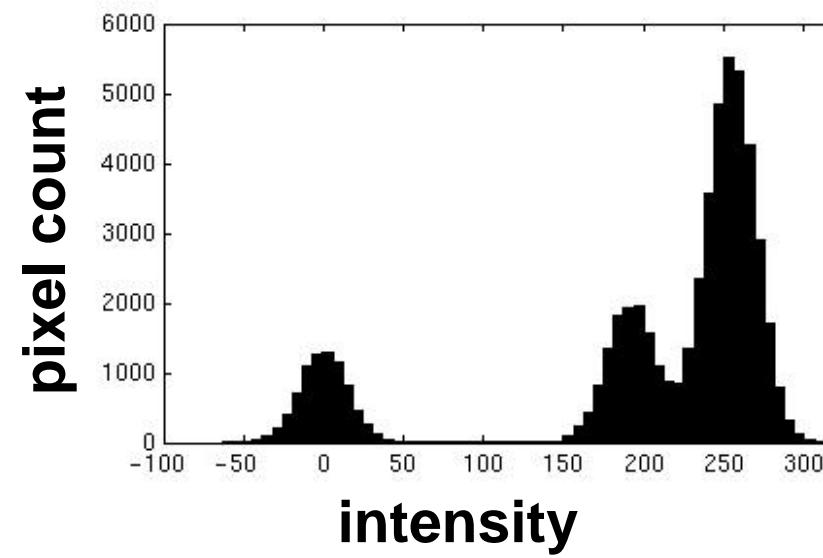
Noisy Images



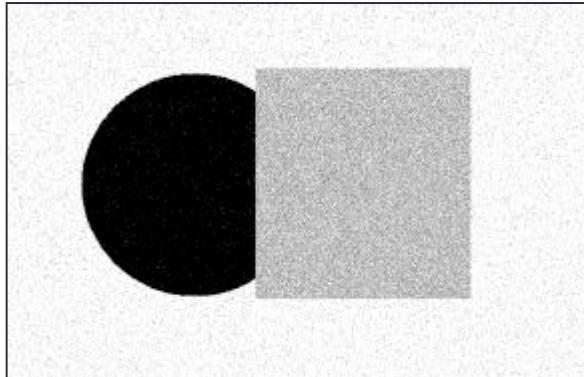
input image



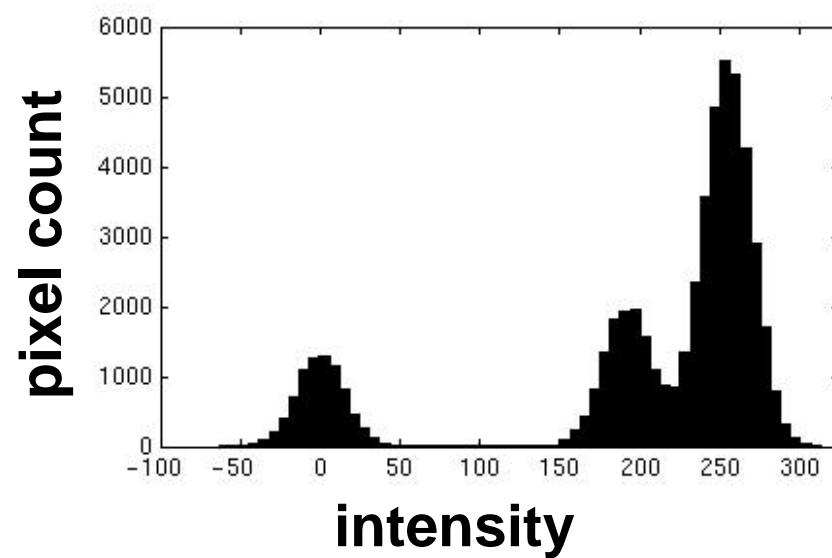
input image



Noisy Images

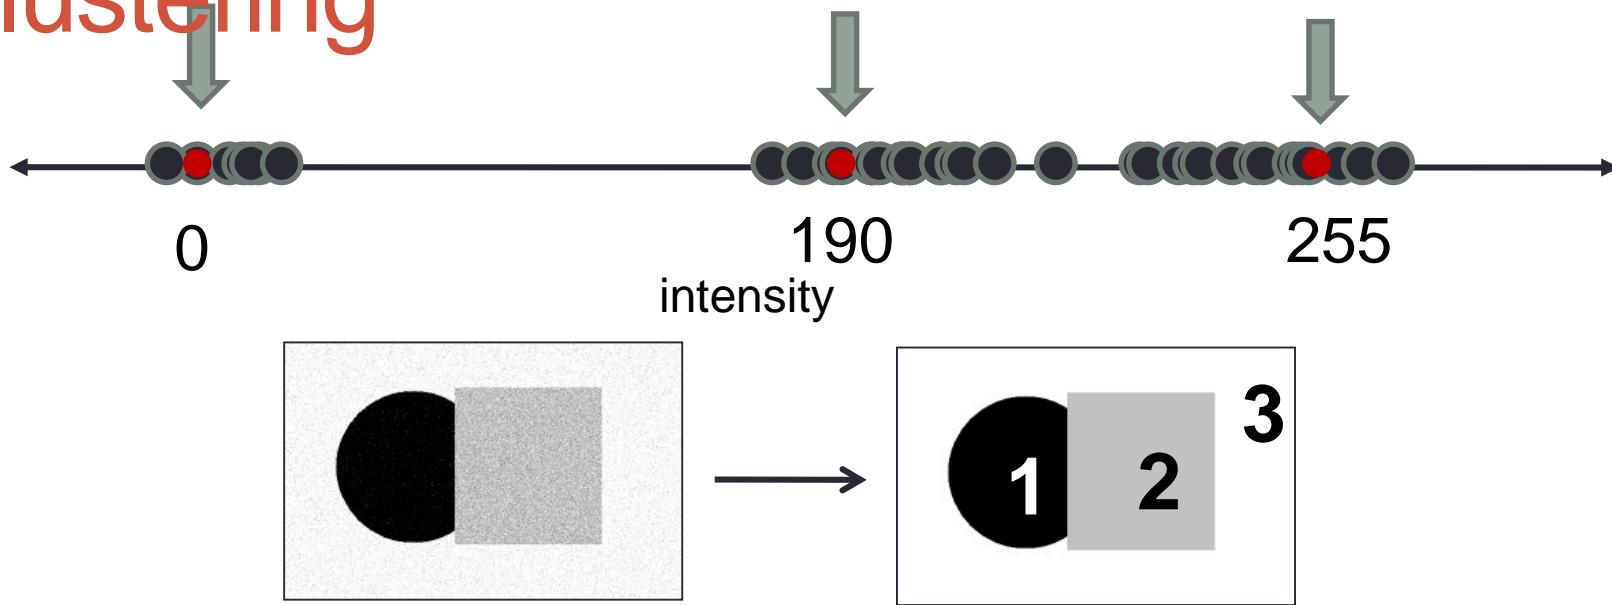


input image



- Now how to determine the three main intensities that define our groups?
- We need to ***cluster***.

Clustering

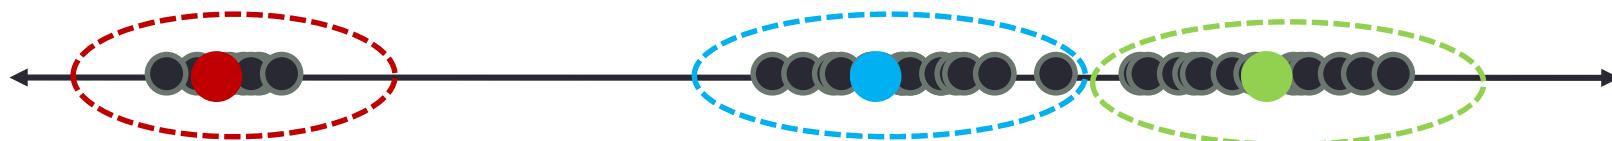


- Goal: choose three “centers” as the **representative** intensities, and label every pixel according to which of these centers it is nearest to.
- Best cluster centers are those that minimize SSD between all points and their nearest cluster center c_i :

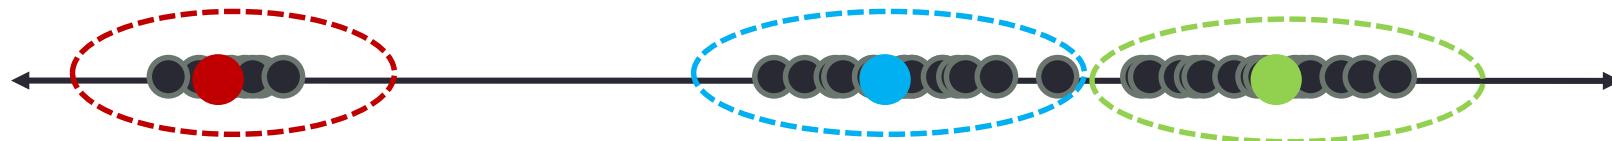
$$\sum_{\text{clusters } i} \sum_{\text{points } p \text{ in cluster } i} \|p - c_i\|^2$$

Clustering

- With this objective, it is a “chicken and egg” problem:
 - Q: how to determine which points to associate with each **cluster center**, c_i ?
 - A: for each point p , choose closest c_i



- Q: If we knew the **group memberships**, how do we get the centers?
- A: choose c_i to be the mean of all points in the cluster

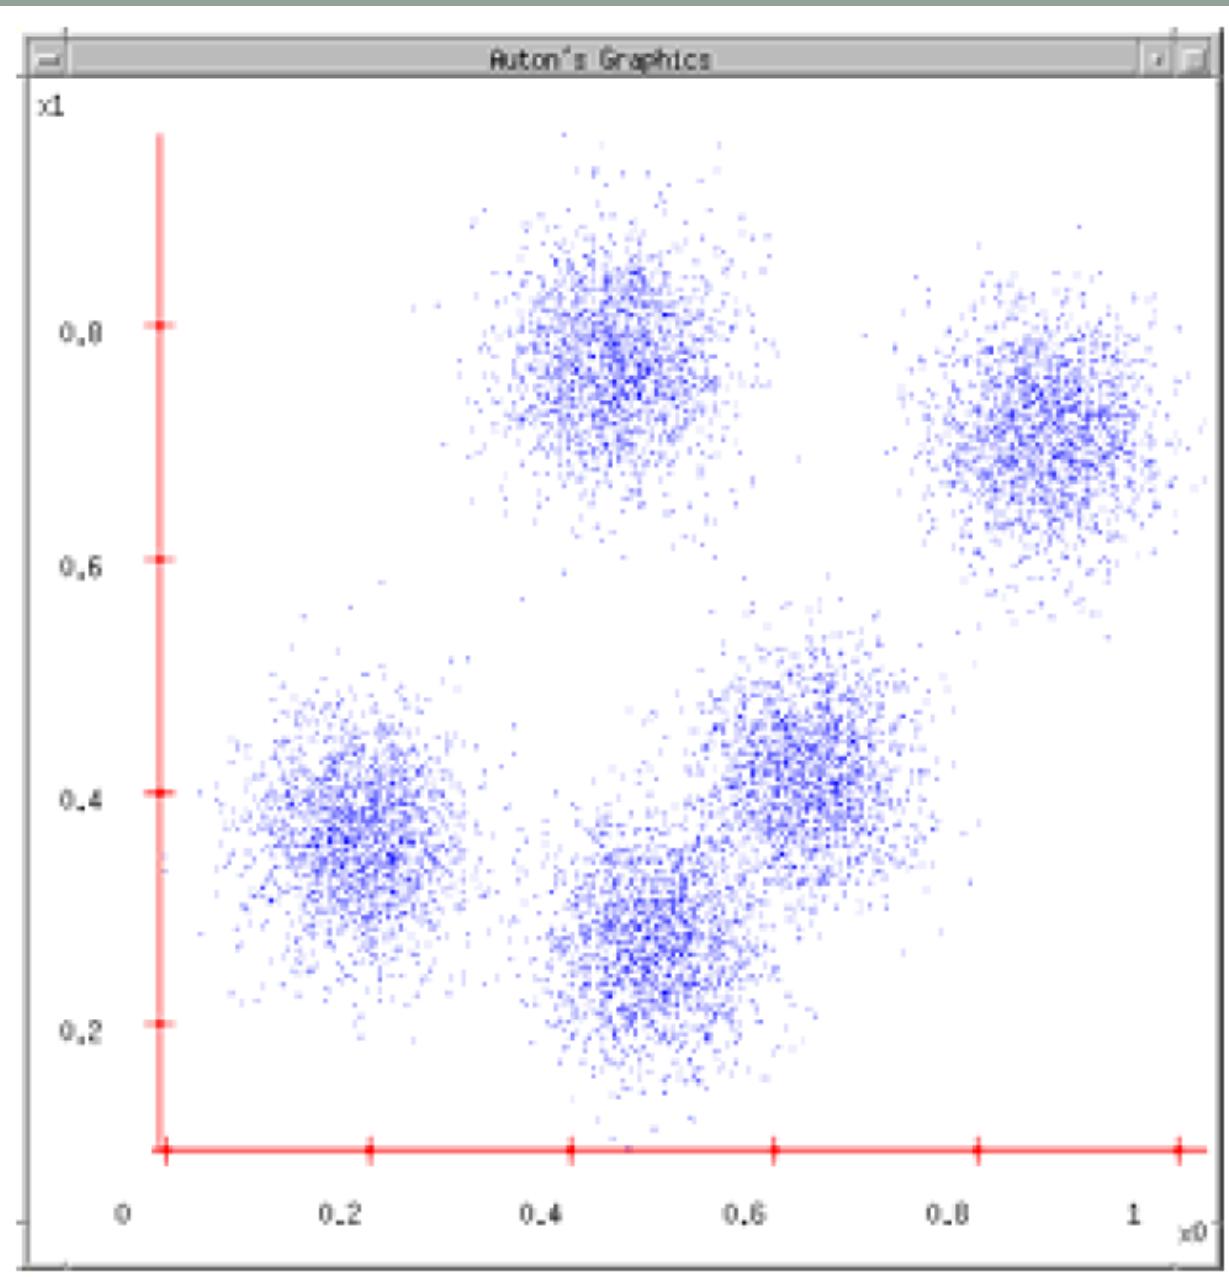


K-means clustering: Algorithm

- 
1. Randomly initialize the cluster centers,
 c_1, \dots, c_K
 2. Given cluster centers, determine points in each cluster
 - For each point p , find the closest c_i . Put p into cluster i
 3. Given points in each cluster, solve for c_i
 - Set c_i to be the mean of points in cluster i
 4. If c_i have changed, repeat Step 2

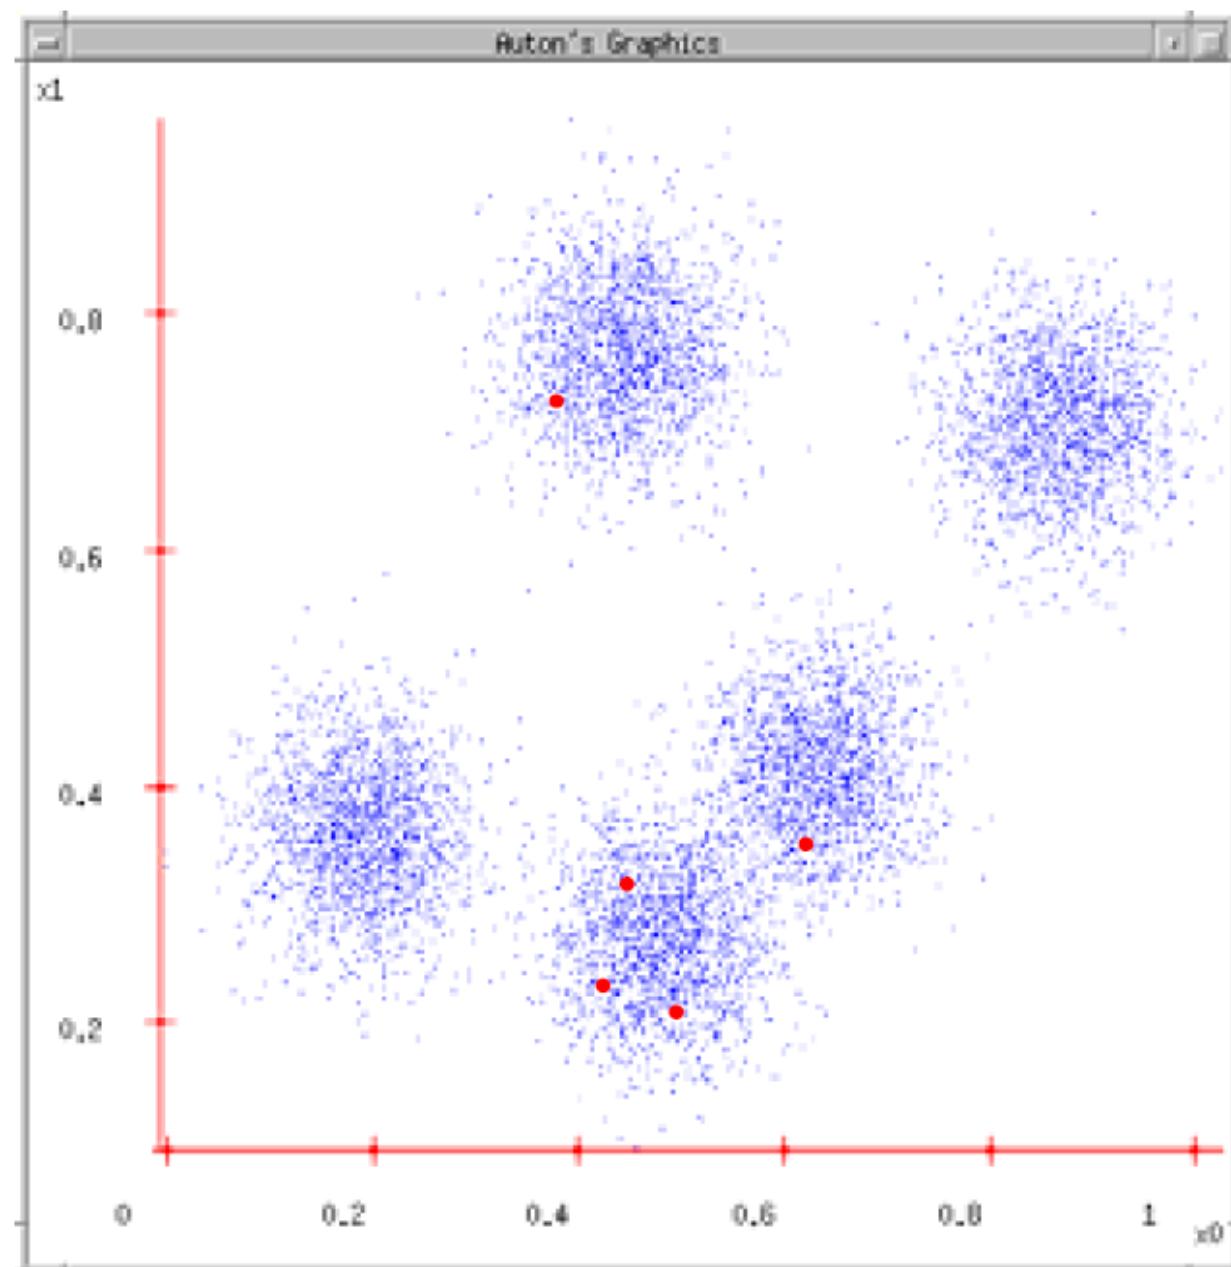
K-means

1. Ask user how many clusters they'd like.
(e.g. $k=5$)



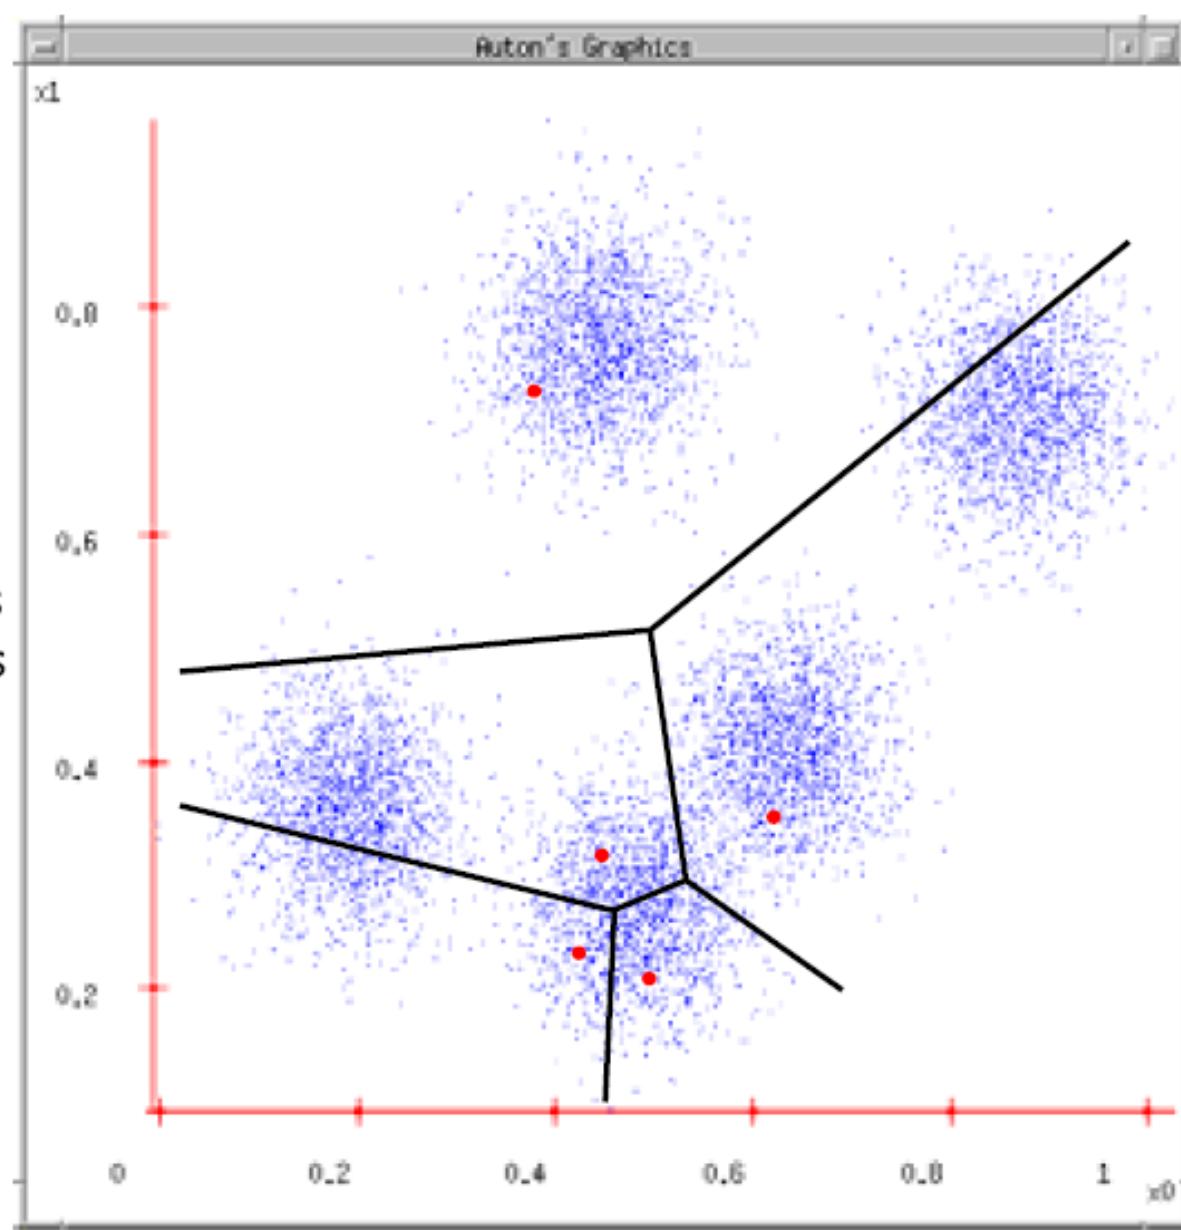
K-means

1. Ask user how many clusters they'd like.
(e.g. $k=5$)
2. Randomly guess k cluster Center locations



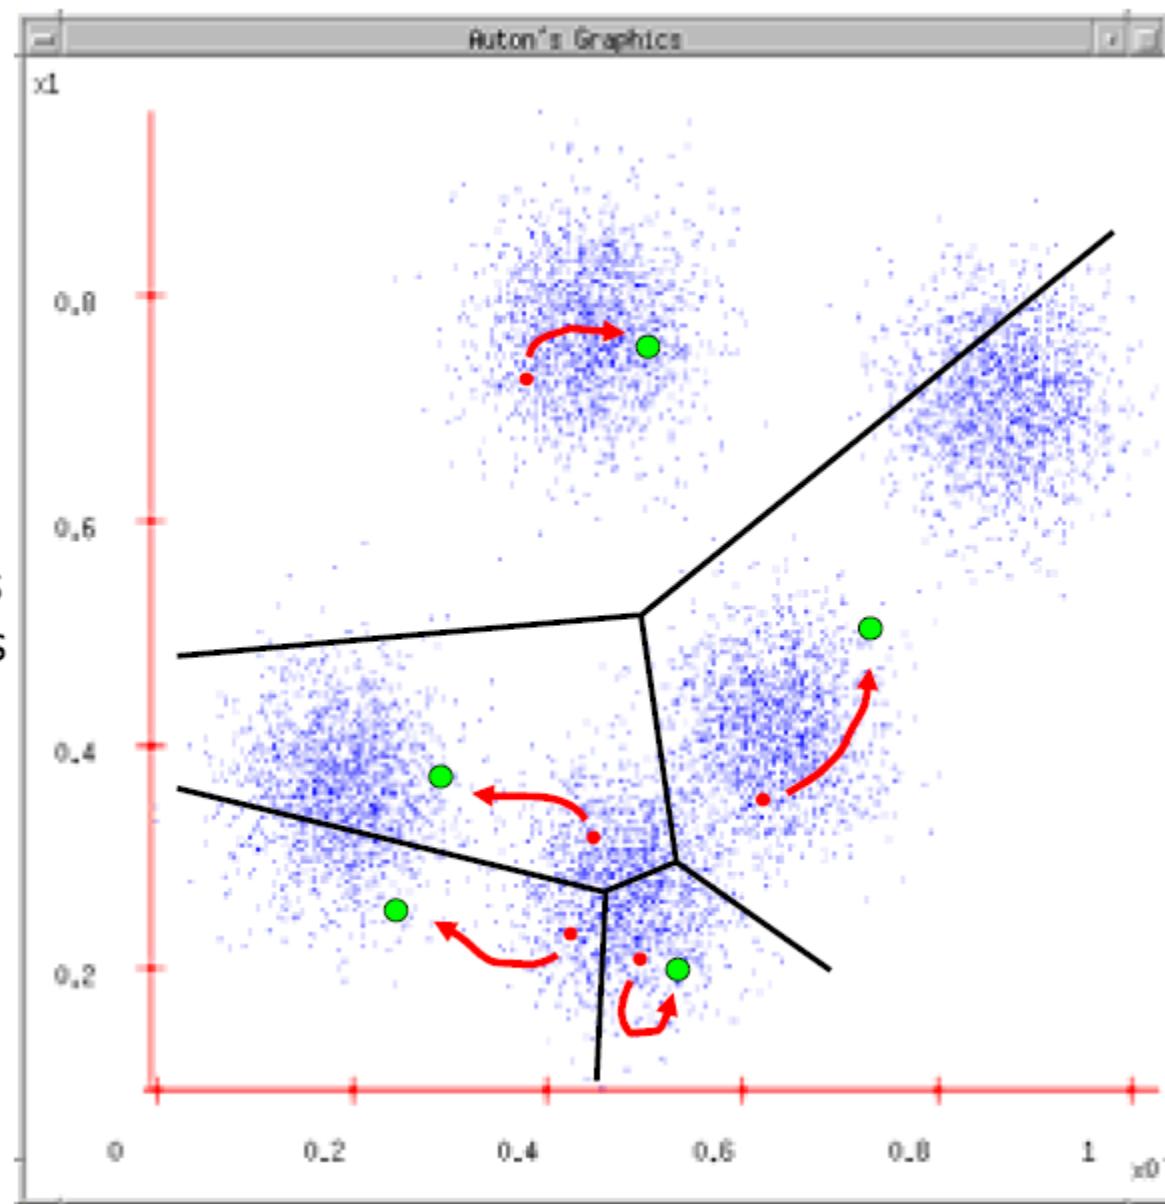
K-means

1. Ask user how many clusters they'd like.
(e.g. $k=5$)
2. Randomly guess k cluster Center locations
3. Each datapoint finds out which Center it's closest to. (Thus each Center "owns" a set of datapoints)



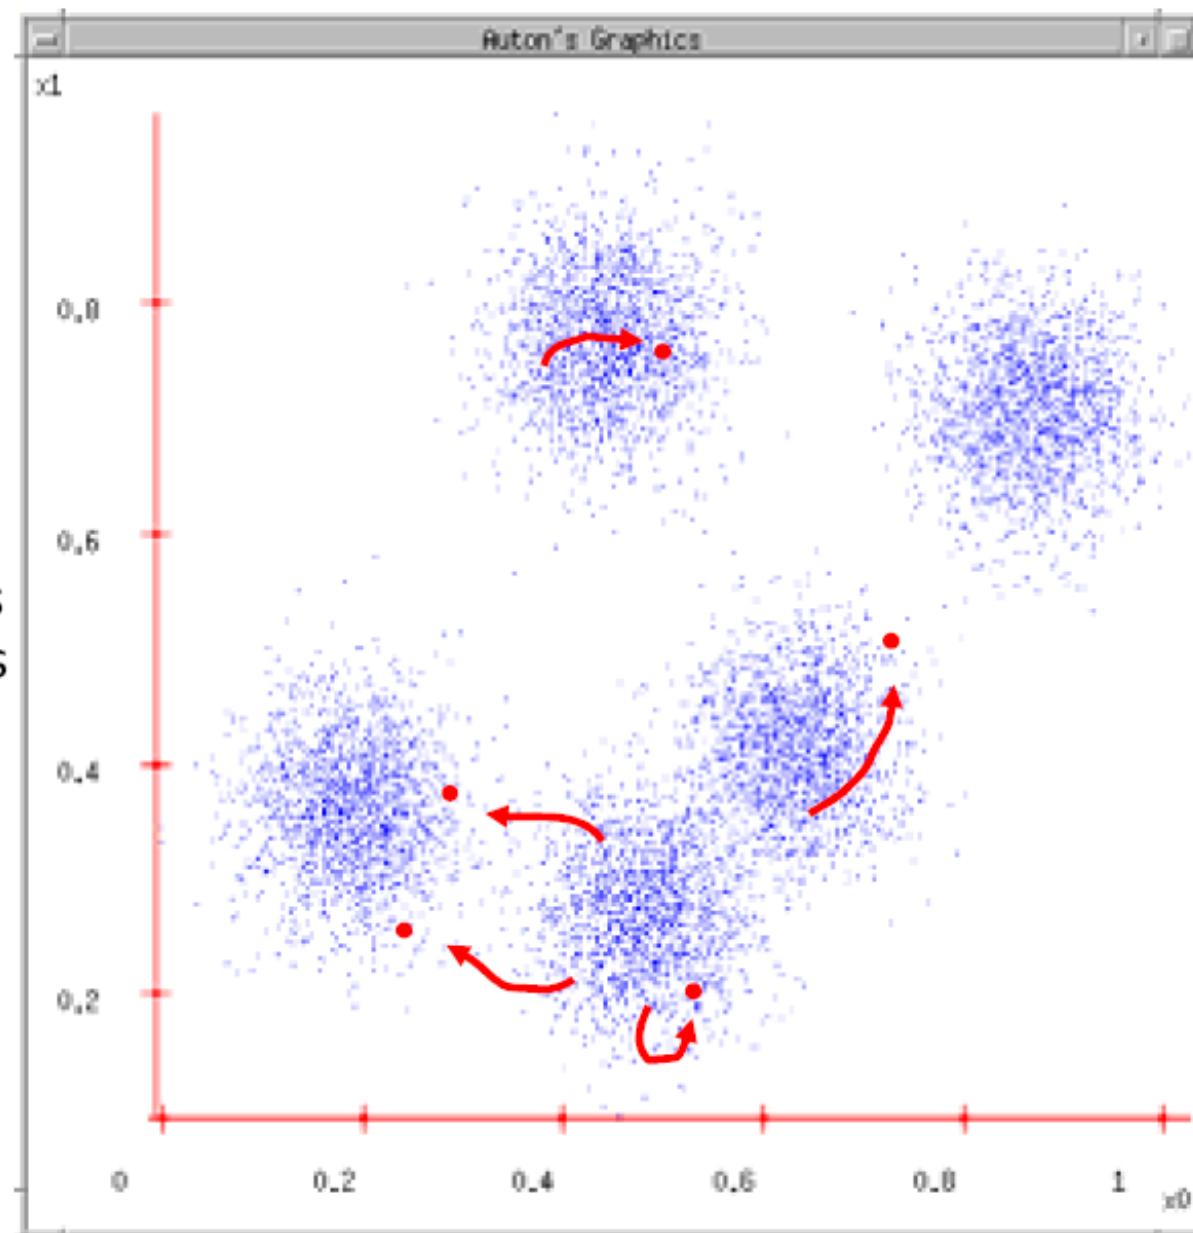
K-means

1. Ask user how many clusters they'd like.
(e.g. $k=5$)
2. Randomly guess k cluster Center locations
3. Each datapoint finds out which Center it's closest to.
4. Each Center finds the centroid of the points it owns



K-means

1. Ask user how many clusters they'd like.
(e.g. $k=5$)
2. Randomly guess k cluster Center locations
3. Each datapoint finds out which Center it's closest to.
4. Each Center finds the centroid of the points it owns...
5. ...and jumps there
6. ...Repeat until terminated!



Segmentation as clustering

Depending on what we choose as the *feature space*, we can group pixels in different ways.

Grouping pixels based
on **intensity** similarity



Feature space: intensity value (1-d)

Number of Clusters



$K=2$



$K=3$

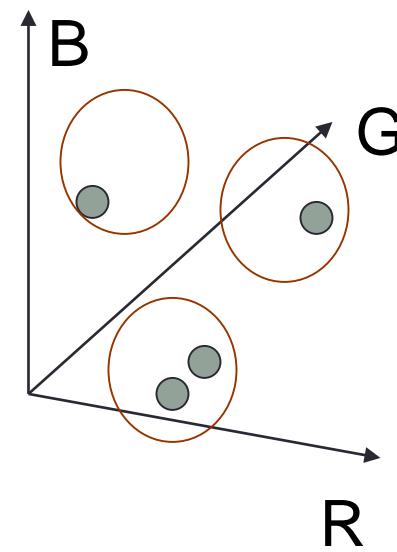


*quantization of the feature space;
segmentation label map*

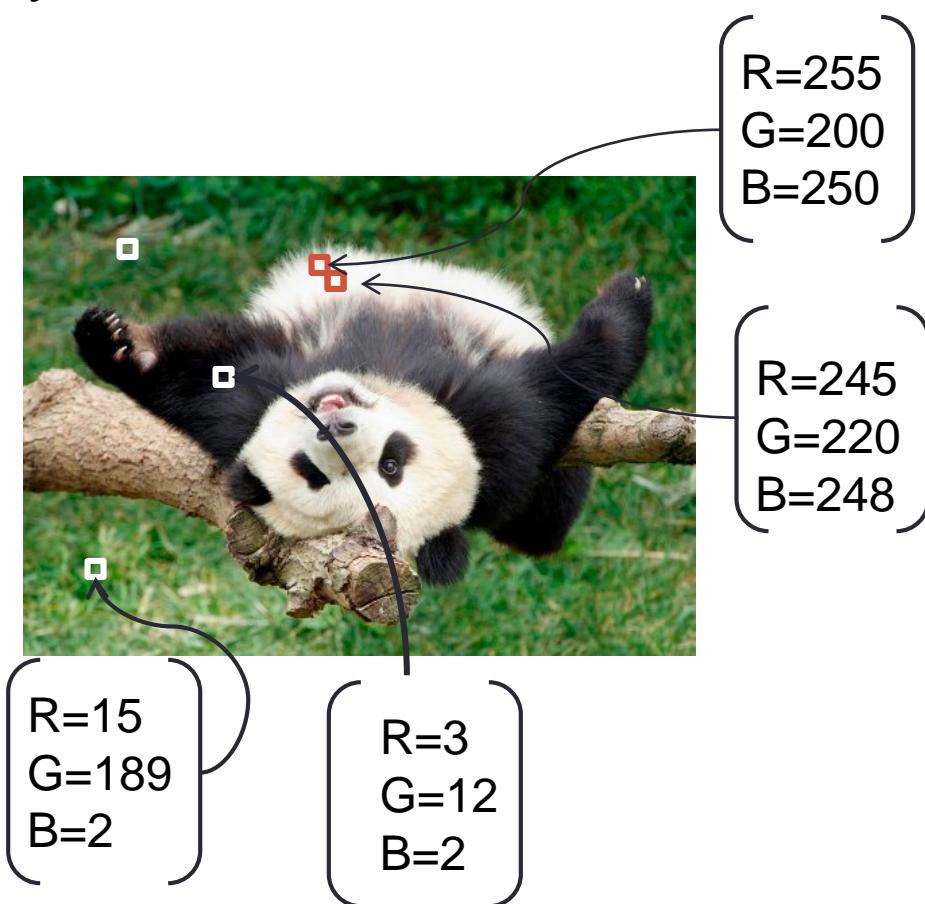
Segmentation as clustering

Depending on what we choose as the *feature space*, we can group pixels in different ways.

Grouping pixels based on **color** similarity

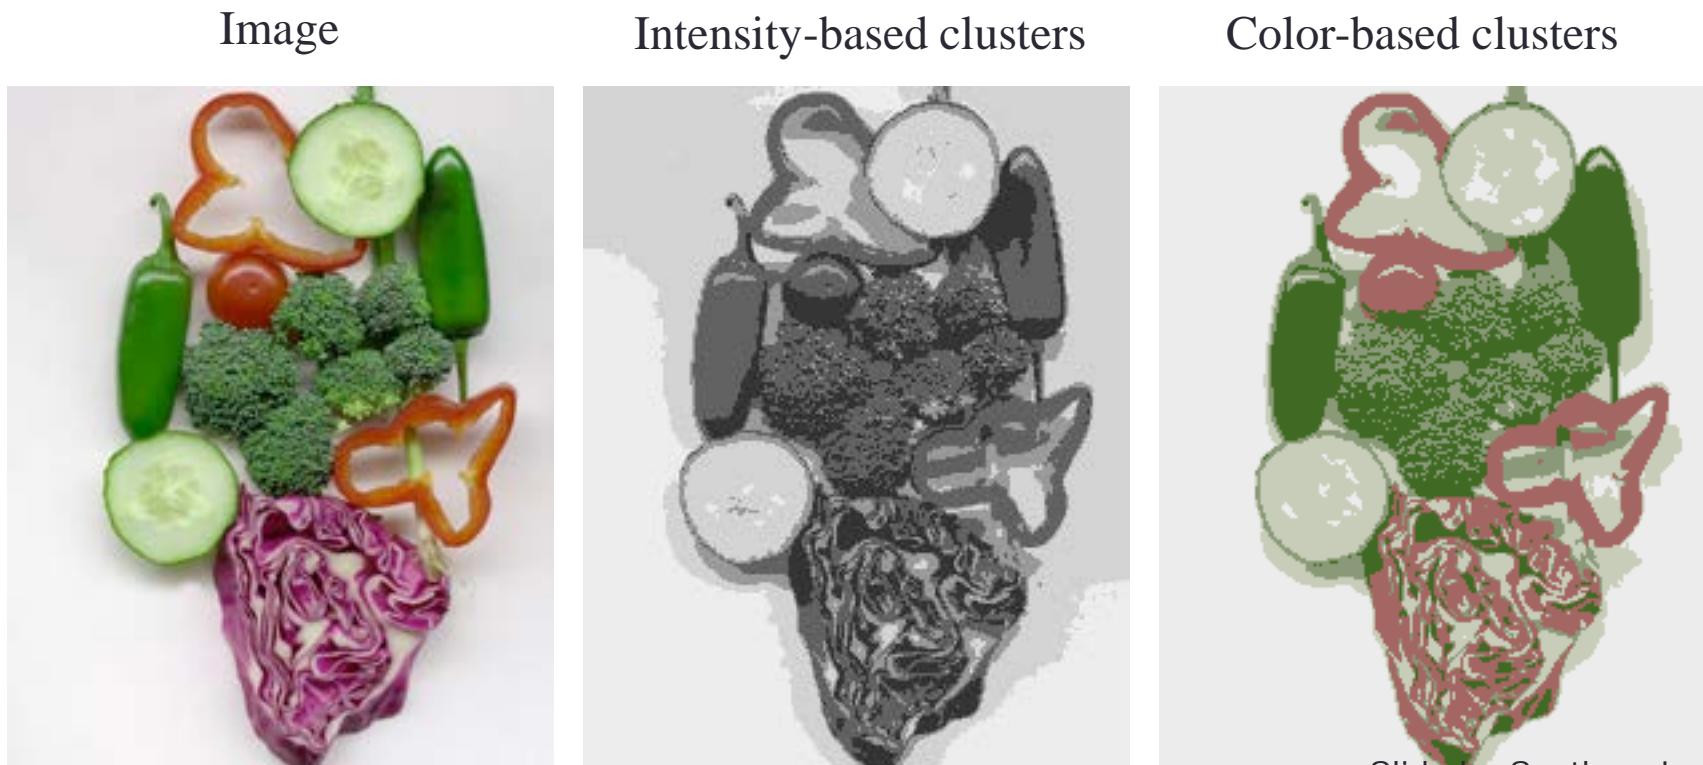


Feature space: color value (3-d)



Segmentation as clustering

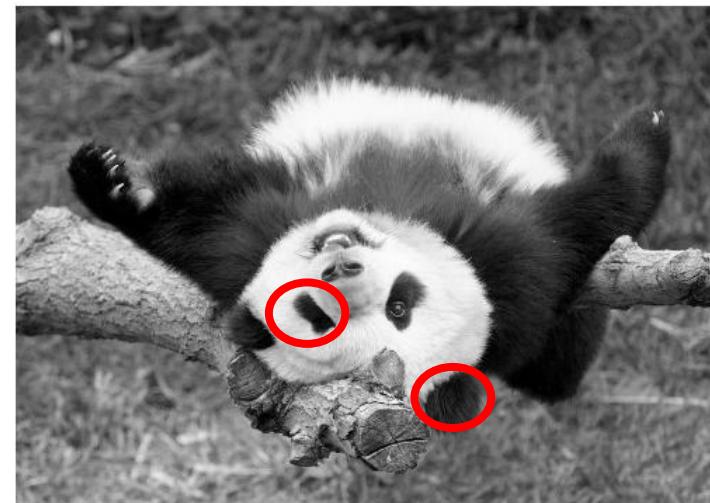
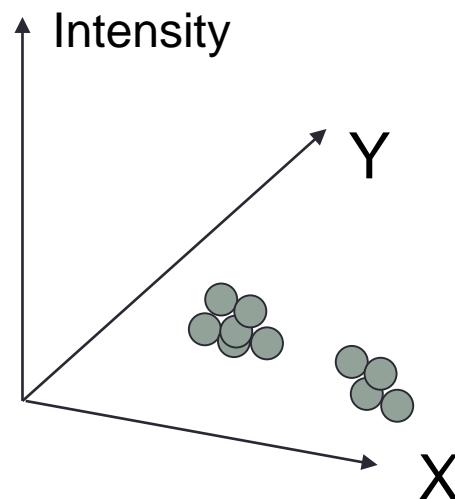
- K-means clustering based on intensity or color is essentially vector quantization of the image attributes



Segmentation as clustering

Depending on what we choose as the *feature space*, we can group pixels in different ways.

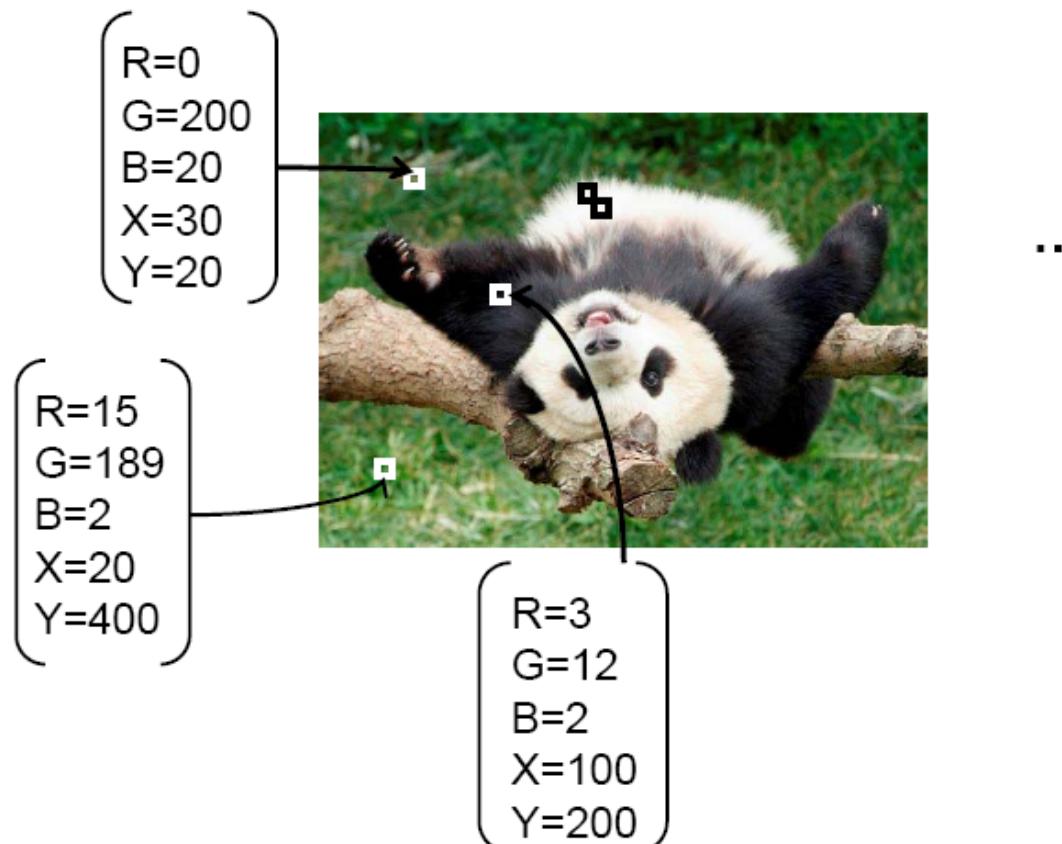
Grouping pixels based on
intensity+position similarity



Both regions are black, but if we also include **position (x,y)**, then we could group the two into distinct segments; way to encode both similarity & proximity.

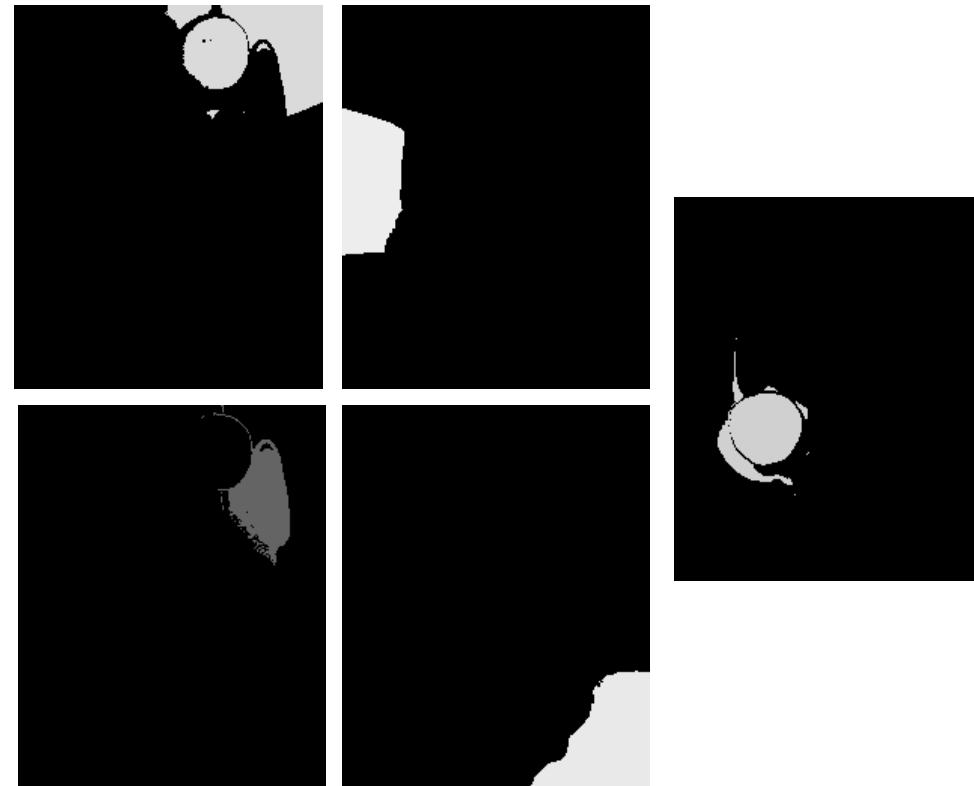
Segmentation as clustering

- Cluster similar pixels (features) together



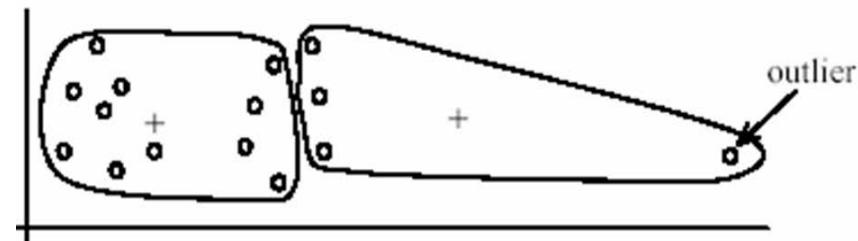
Segmentation as clustering

- Clustering based on (r,g,b,x,y) values enforces more spatial coherence

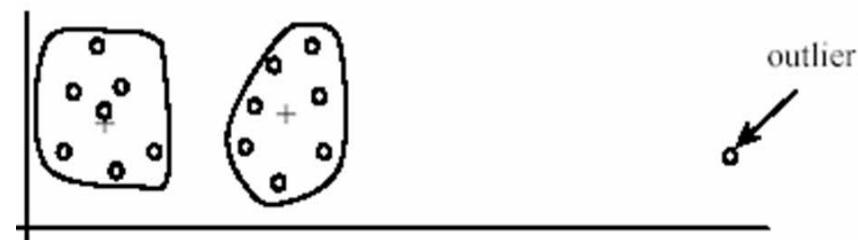


K-Means for segmentation

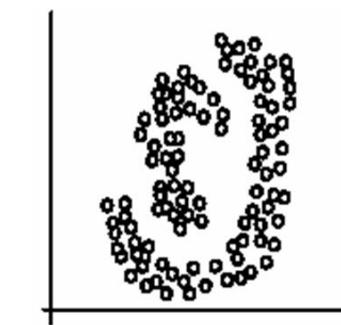
- Pros
 - Very simple method
 - Converges to a local minimum of the error function
- Cons
 - Memory-intensive
 - Need to pick K
 - Sensitive to initialization
 - Sensitive to outliers
 - Only finds “spherical” clusters



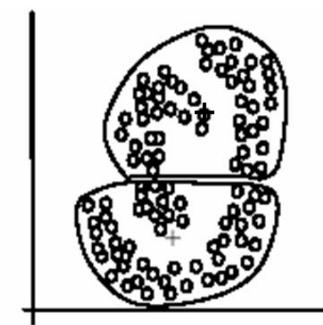
(A): Undesirable clusters



(B): Ideal clusters



(A): Two natural clusters

(B): k -means clusters

Segmentation as clustering

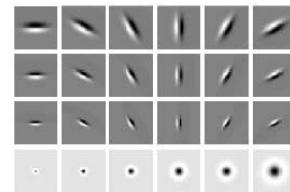
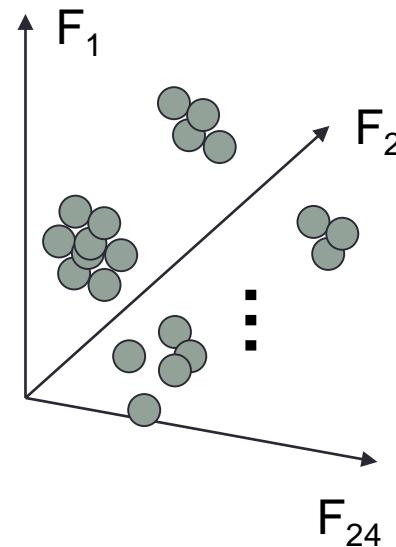
- Color, brightness, position alone are not enough to distinguish all regions...



Segmentation as clustering

Depending on what we choose as the *feature space*, we can group pixels in different ways.

Grouping pixels based on **texture** similarity

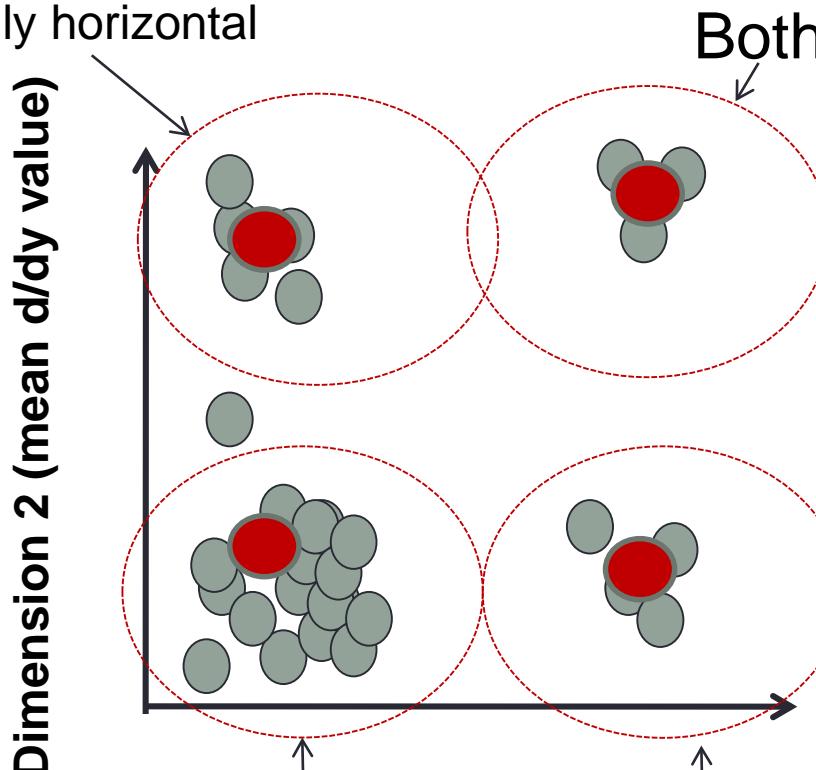


Filter bank
of 24 filters

Feature space: filter bank responses (e.g., 24-d)

Aside: Texture representation example

Windows with primarily horizontal edges



Windows with small gradient in both directions

Windows with primarily vertical edges

Both

	<u>mean d/dx value</u>	<u>mean d/dy value</u>
Win. #1	4	10
Win. #2	18	7
⋮	⋮	⋮
Win. #9	20	20
⋮	⋮	⋮

statistics to summarize patterns in small windows

Aside: Texture features

- Find “textons” by **clustering** vectors of filter bank outputs
- Describe texture in a window based on its *texton histogram*

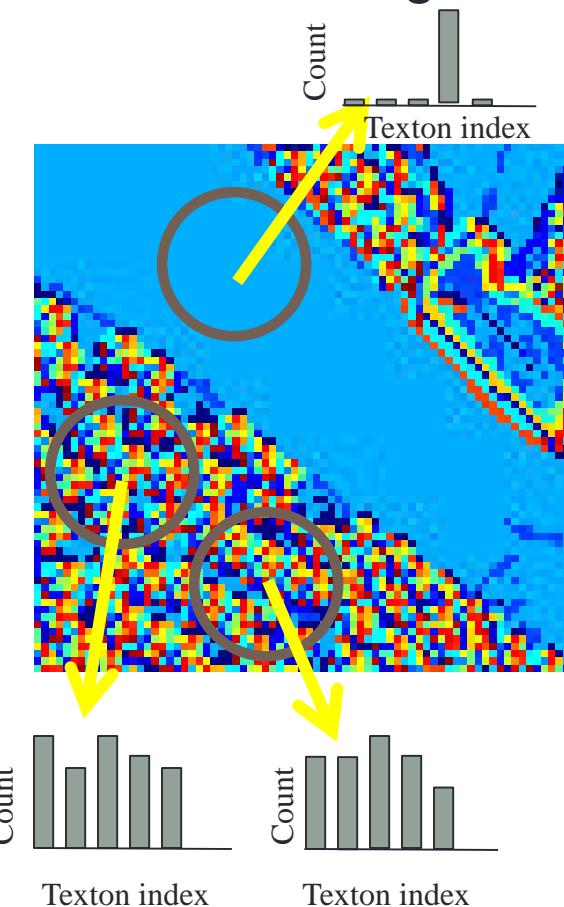
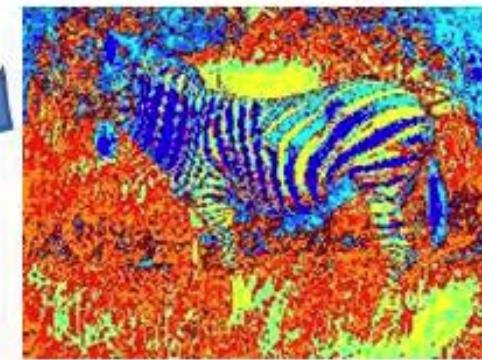
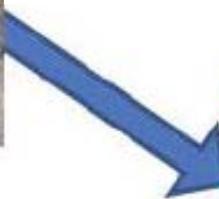


Image segmentation example



Color-based regions

Make it better

- K -means heavily sensitive to initial conditions and (typically) need to know K in advance.
- Suppose we assume that there are a few *modes* in the image and that all the pixels come from these modes.
- If you could find the modes you might be able to segment the image.

Mean shift algorithm

- The mean shift algorithm seeks *modes* or local maxima of density in the feature space

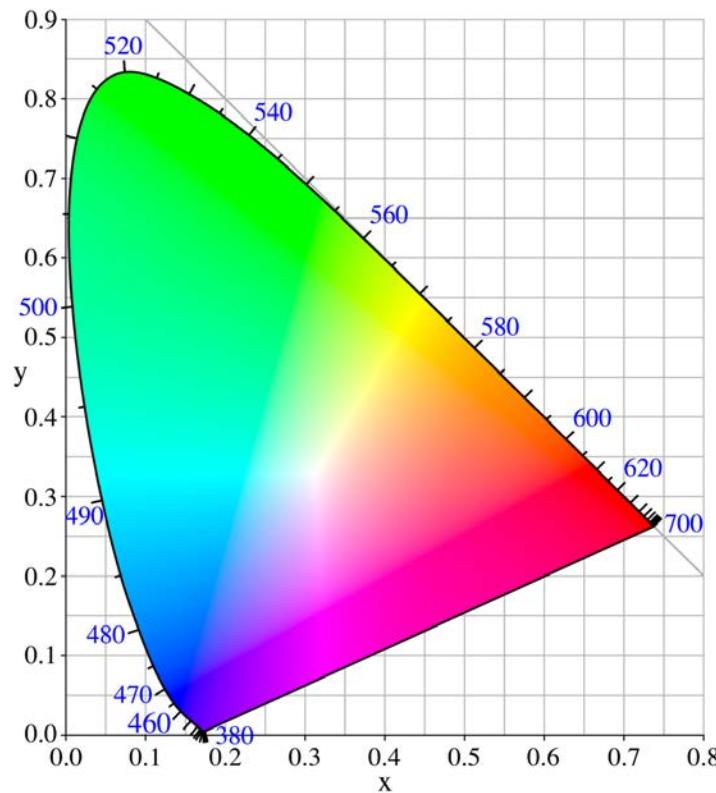
image



A digression about color...

- “Color” is an inherently perceptual phenomena.
- Only related but not the same as wavelength of light energy.
- In fact only some colors are found in the spectrum...

Colors perceivable by the human eye



CIE xy
chromaticity
diagram, 1931

CIE XYZ color space (1931)

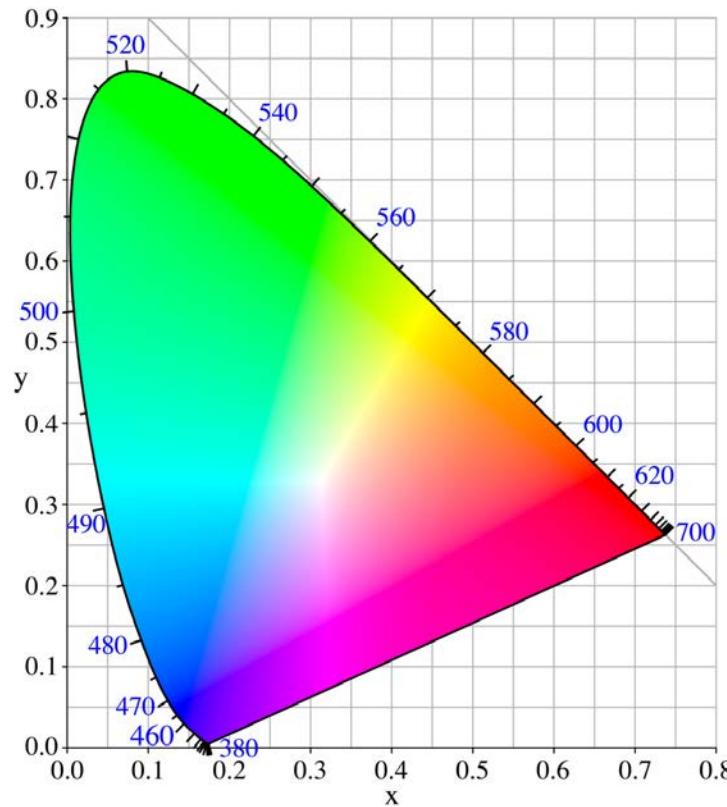
A space with desired properties

- Easy to compute – linear transform of CIE RGB
- Y: Perceived luminance
- X, Z: Perceived color
- Represents a wide range of colors

Colors perceivable by the human eye

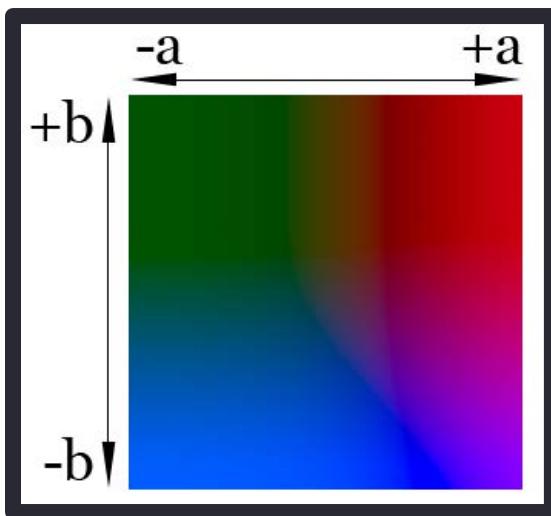
$$y = \frac{X}{X + Y + Z}$$

$$y = \frac{Y}{X + Y + Z}$$

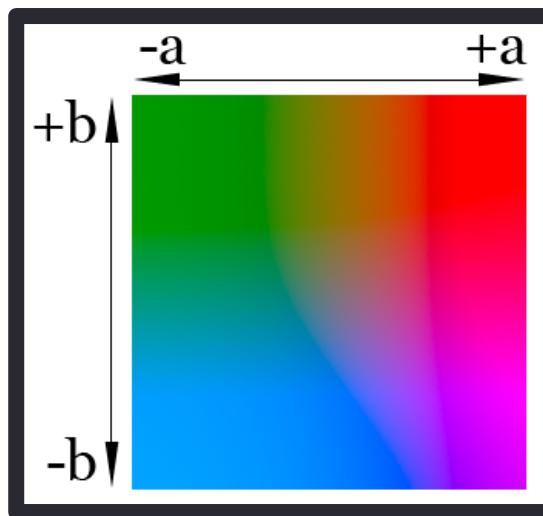


CIE xy
chromaticity
diagram, 1931

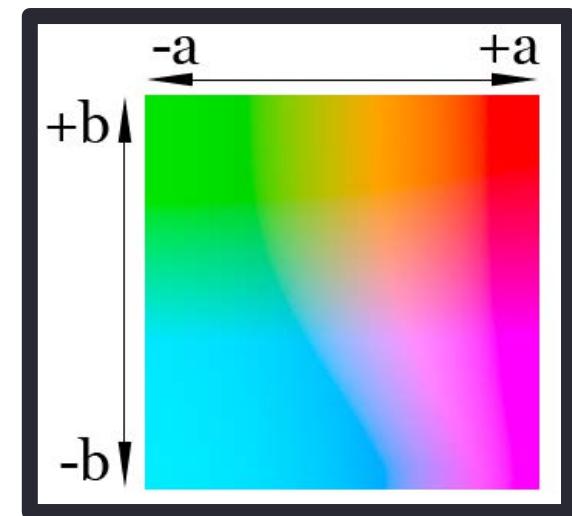
CIE L*a*b* color space



$L = 25\%$



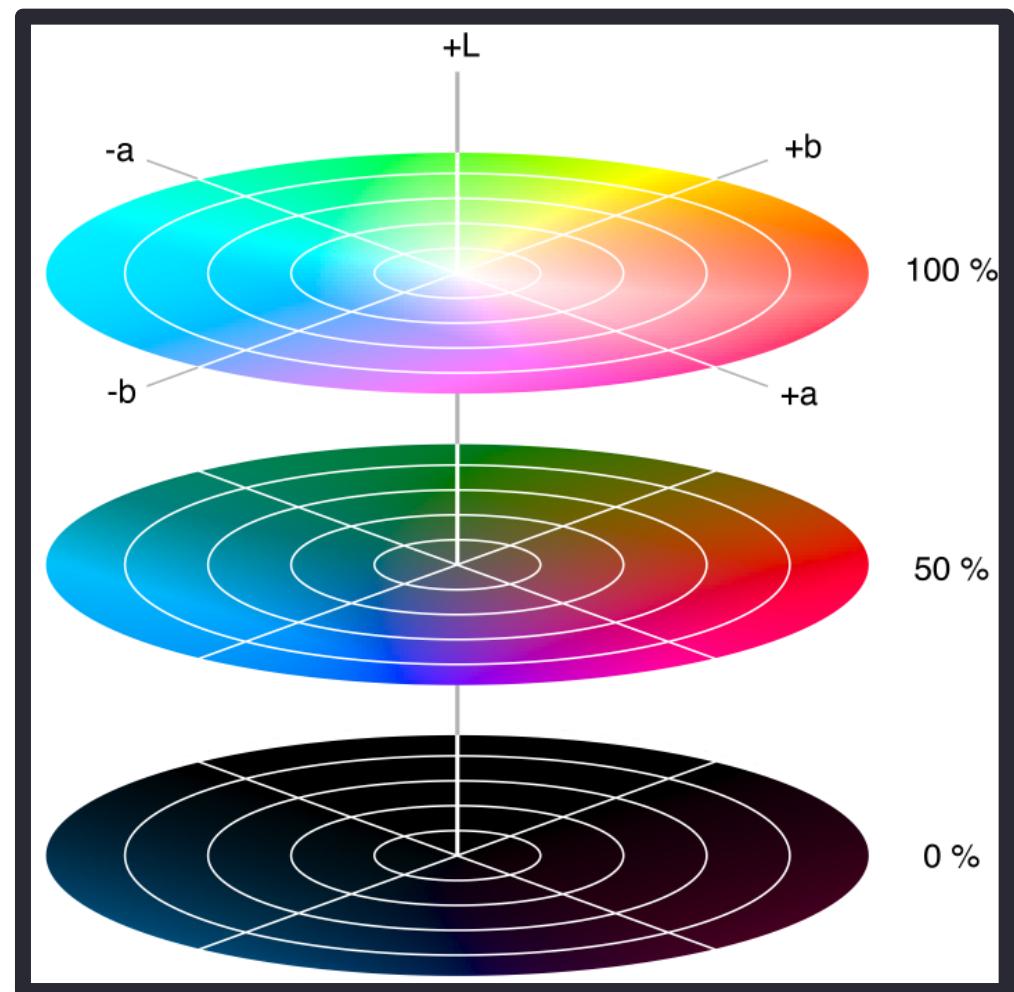
$L = 50\%$



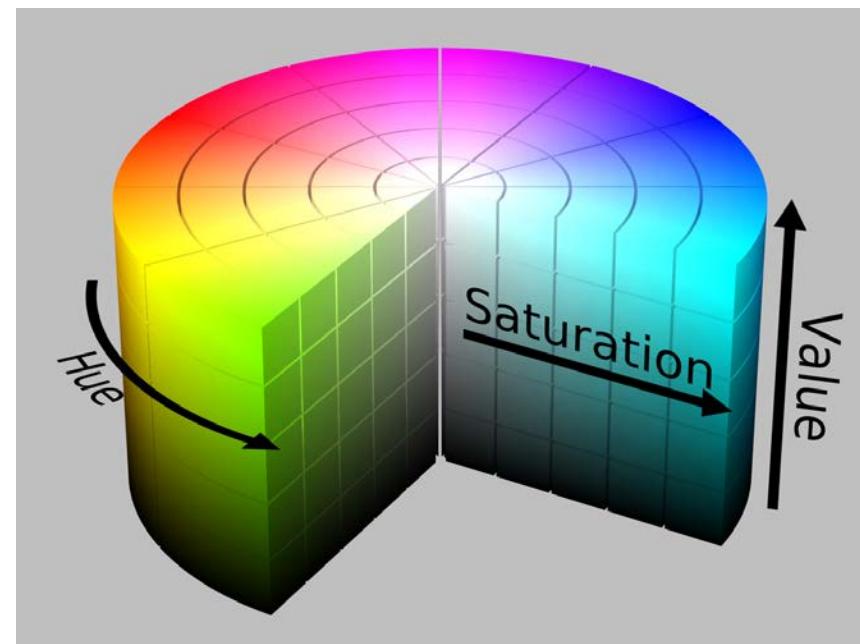
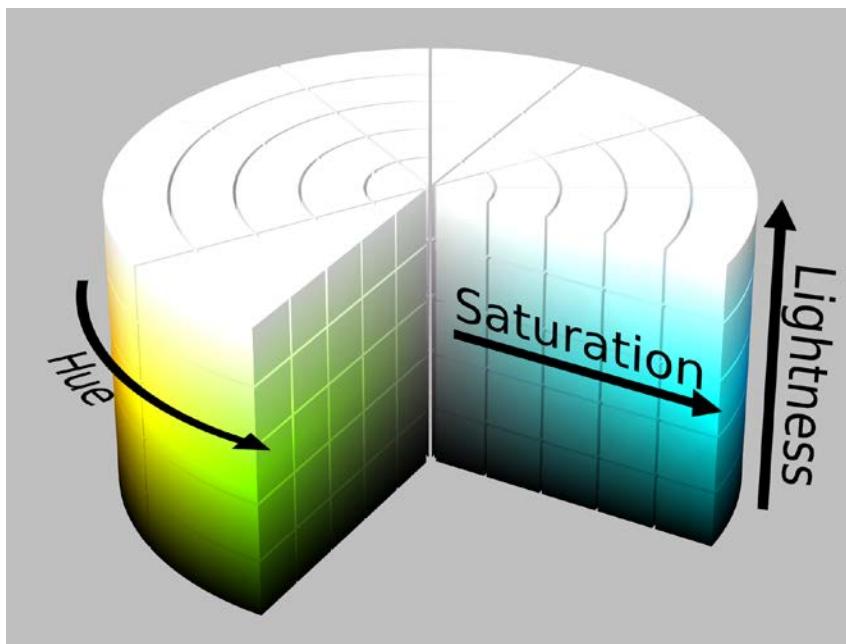
$L = 75\%$

Cylindrical view

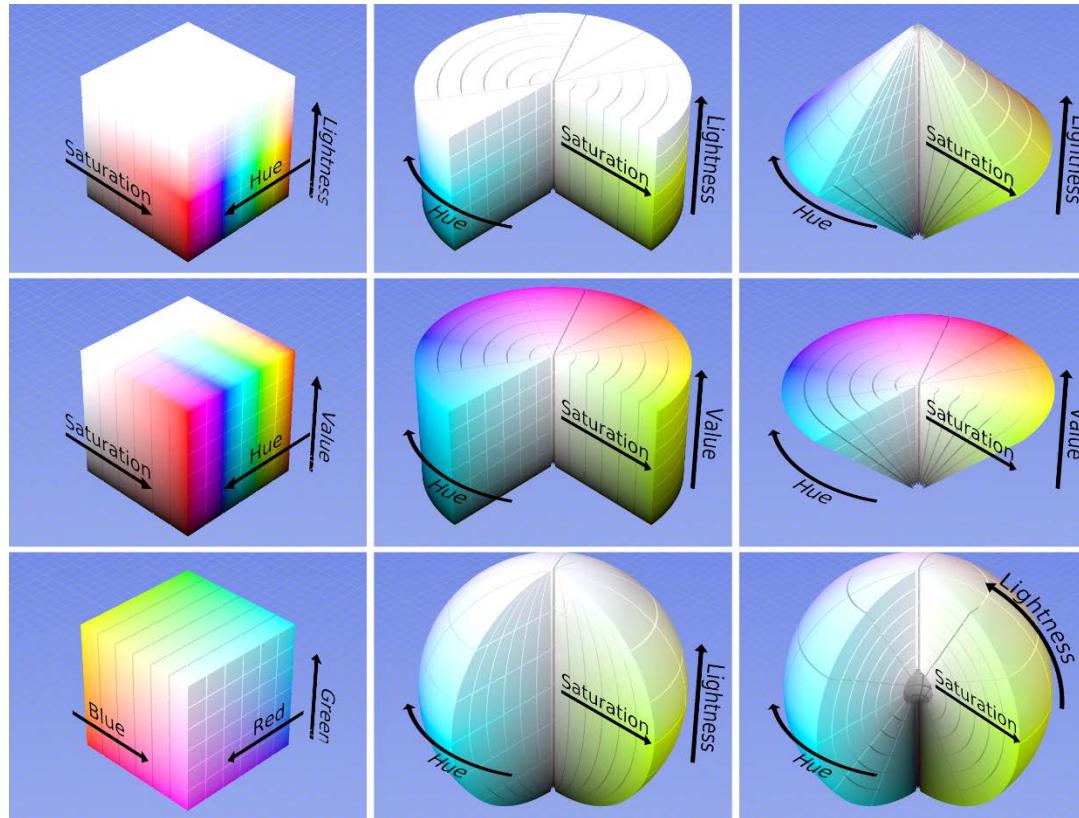
Think of chroma
(here a^* , b^*) defining
a planar disc at each
luminance level (L)



HSL and HSV color spaces

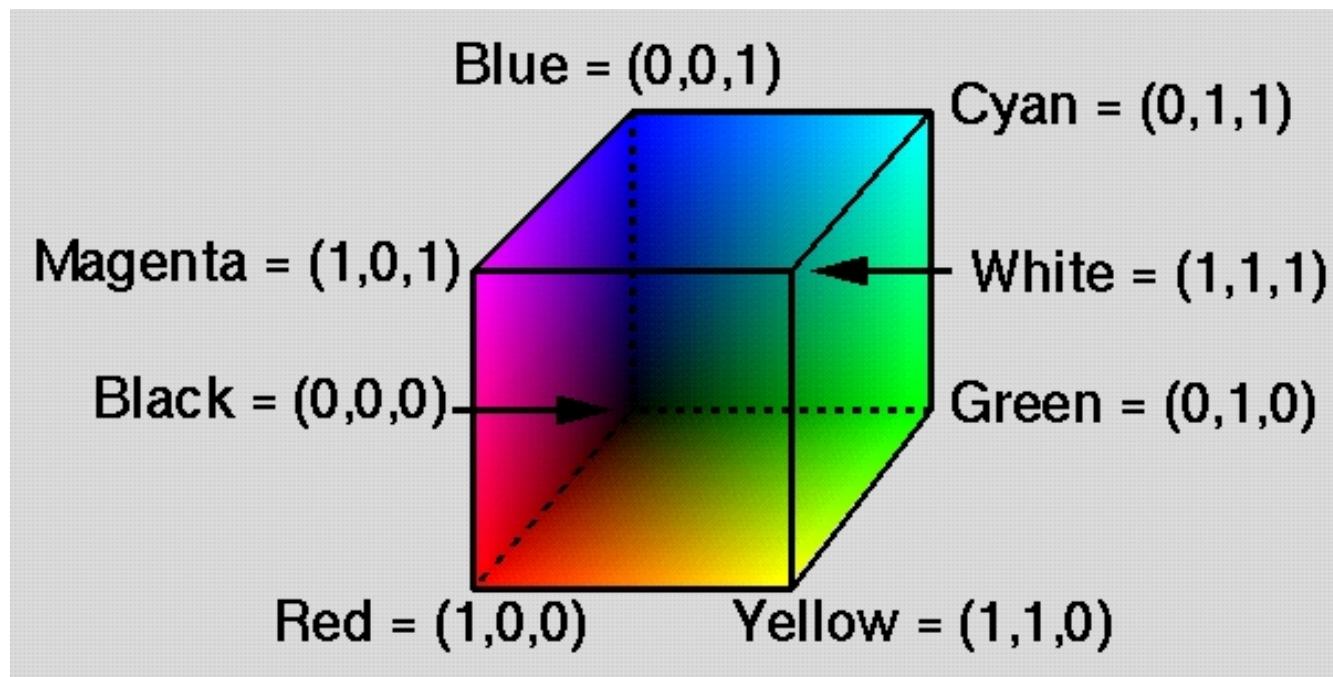


But there are lots of color spaces

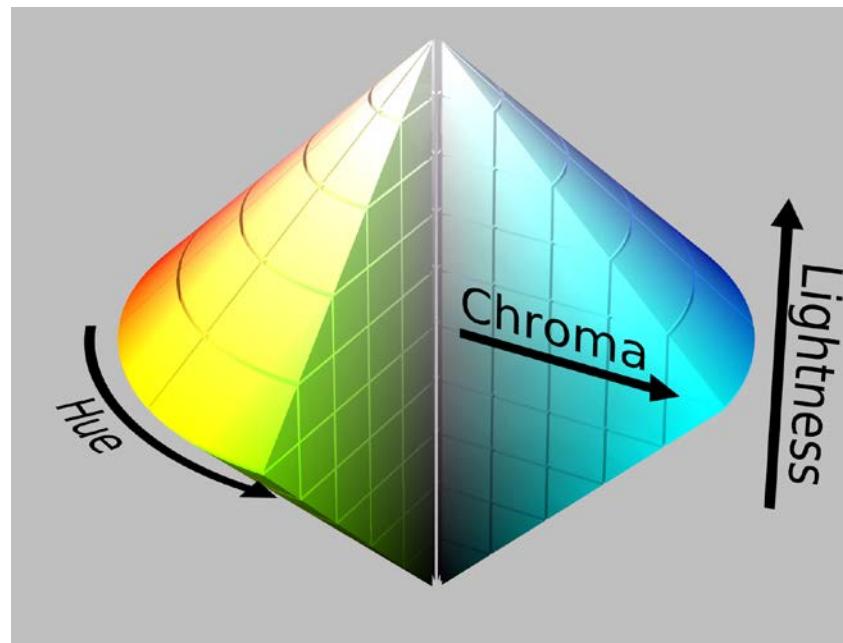


The one we know best...

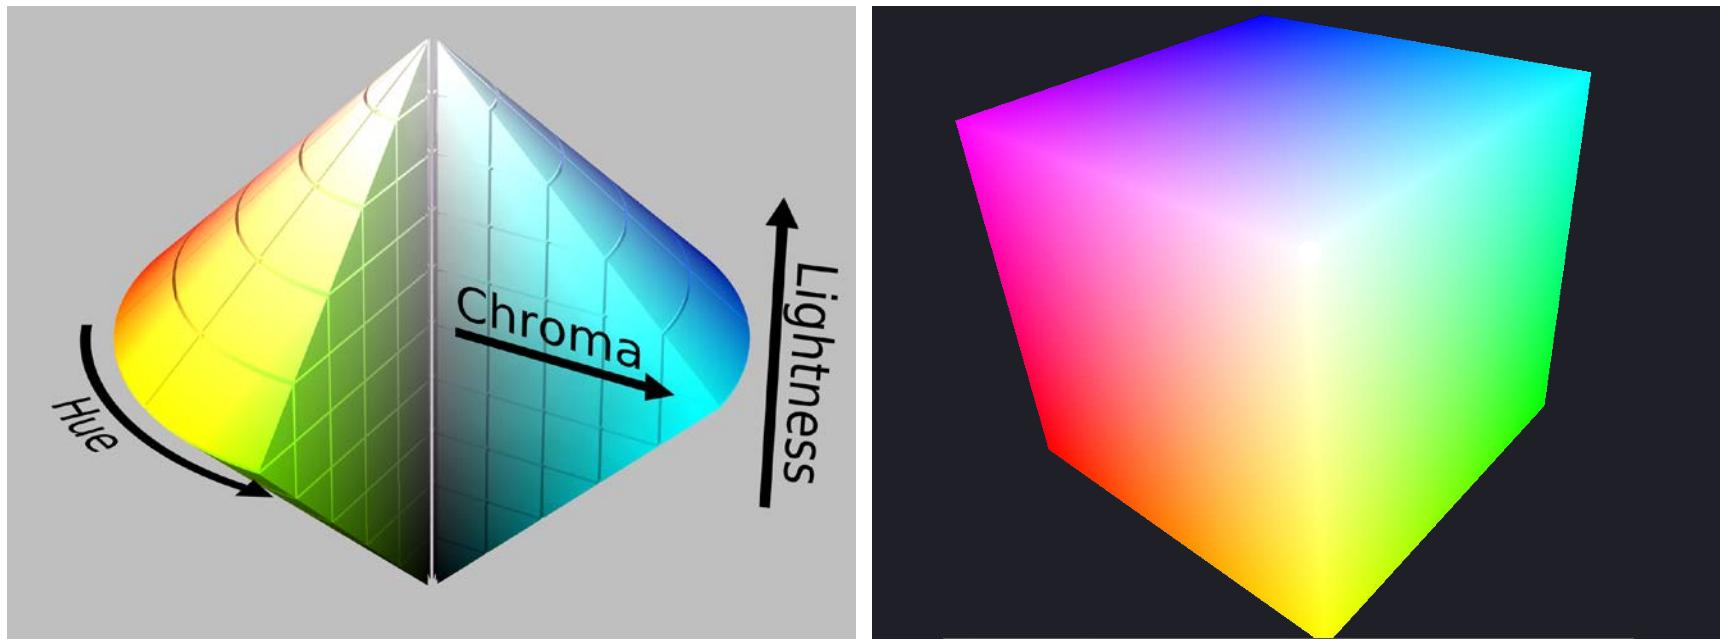
- RGB color space



My favorite



Like a squared double cone?



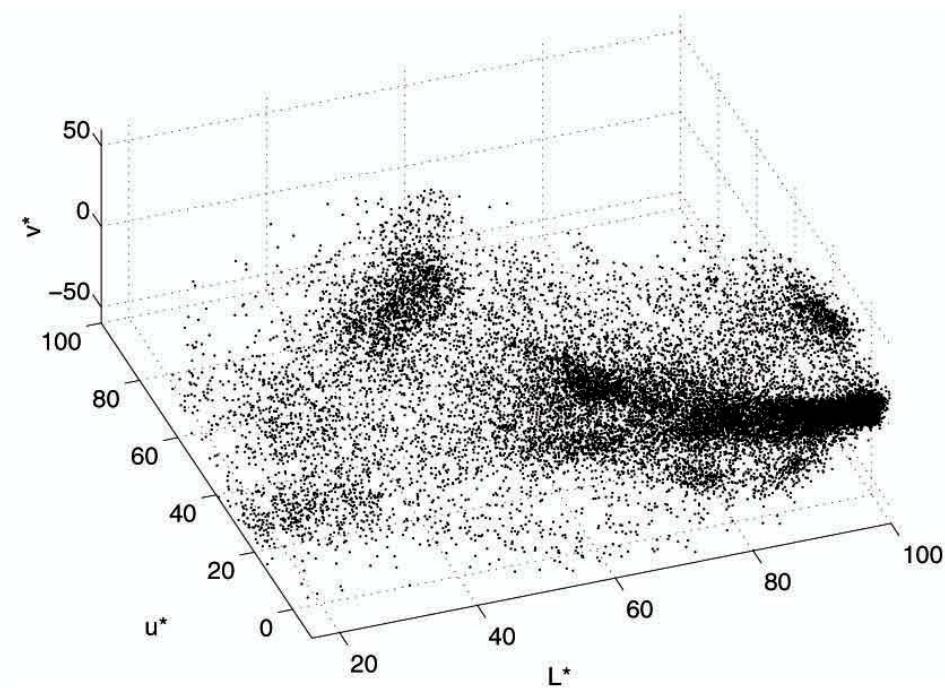
Mean shift algorithm

- The mean shift algorithm seeks *modes* or local maxima of density in the feature space

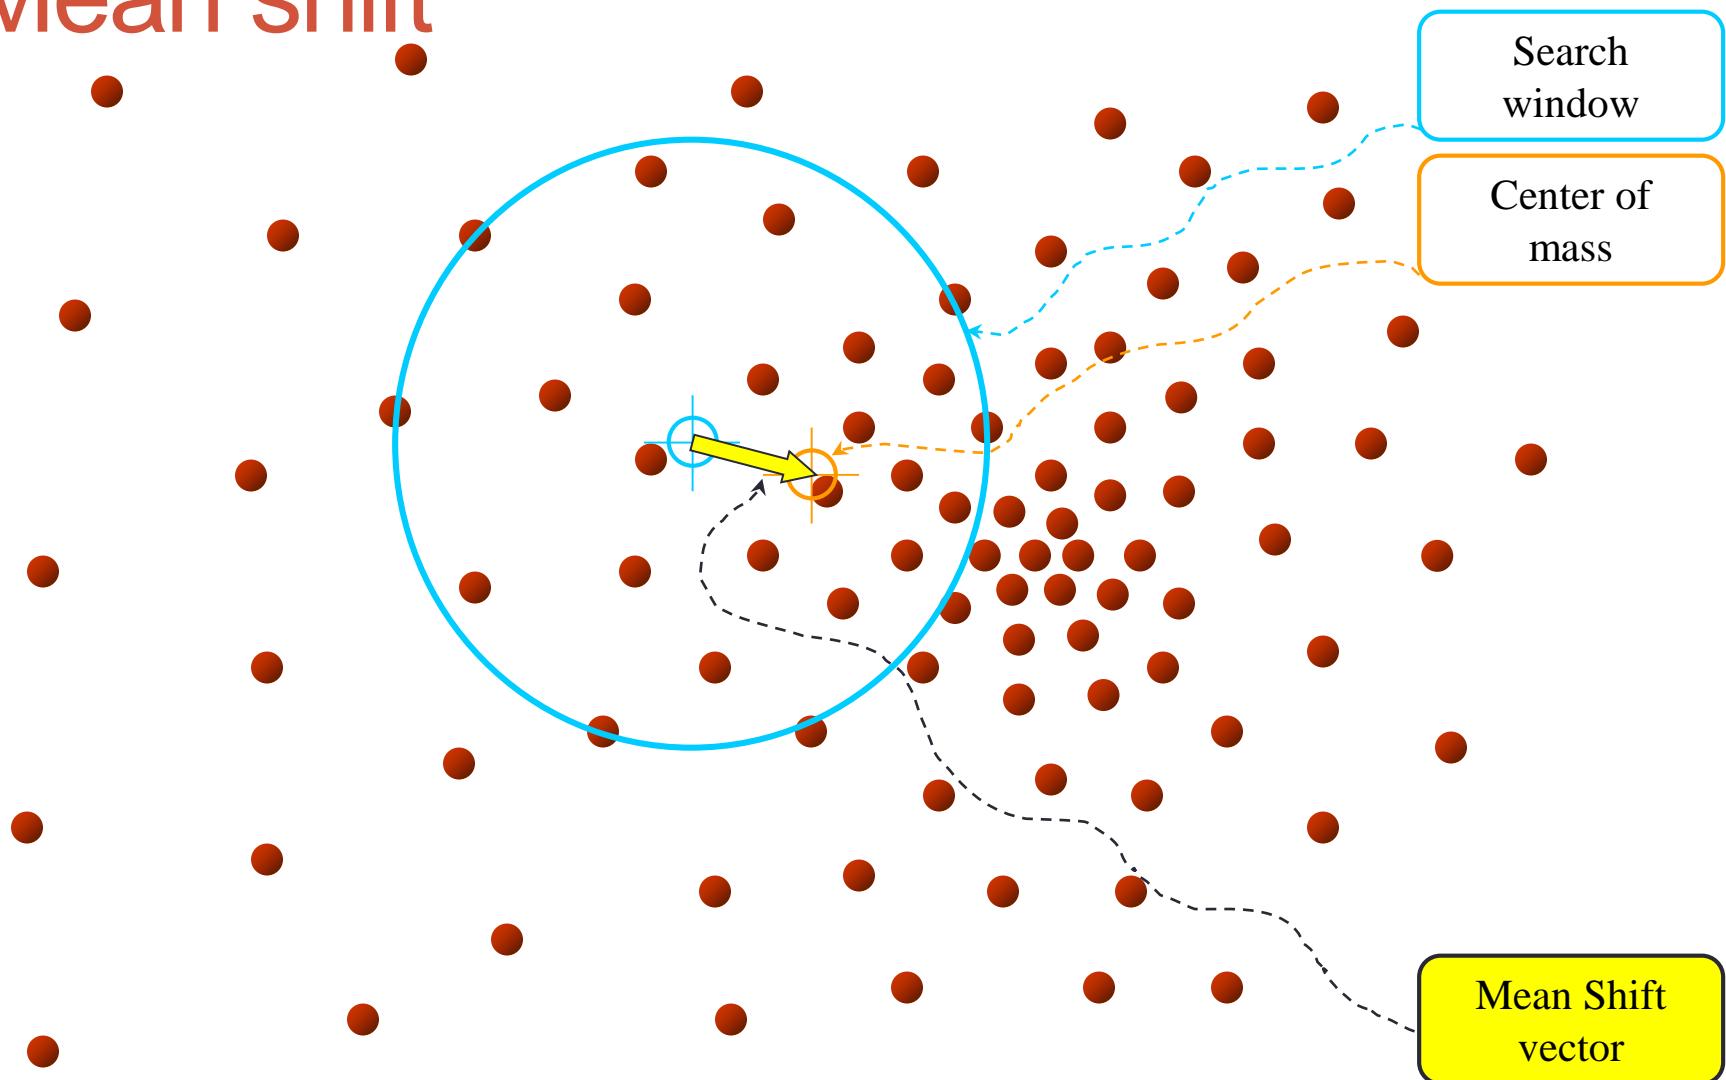
image



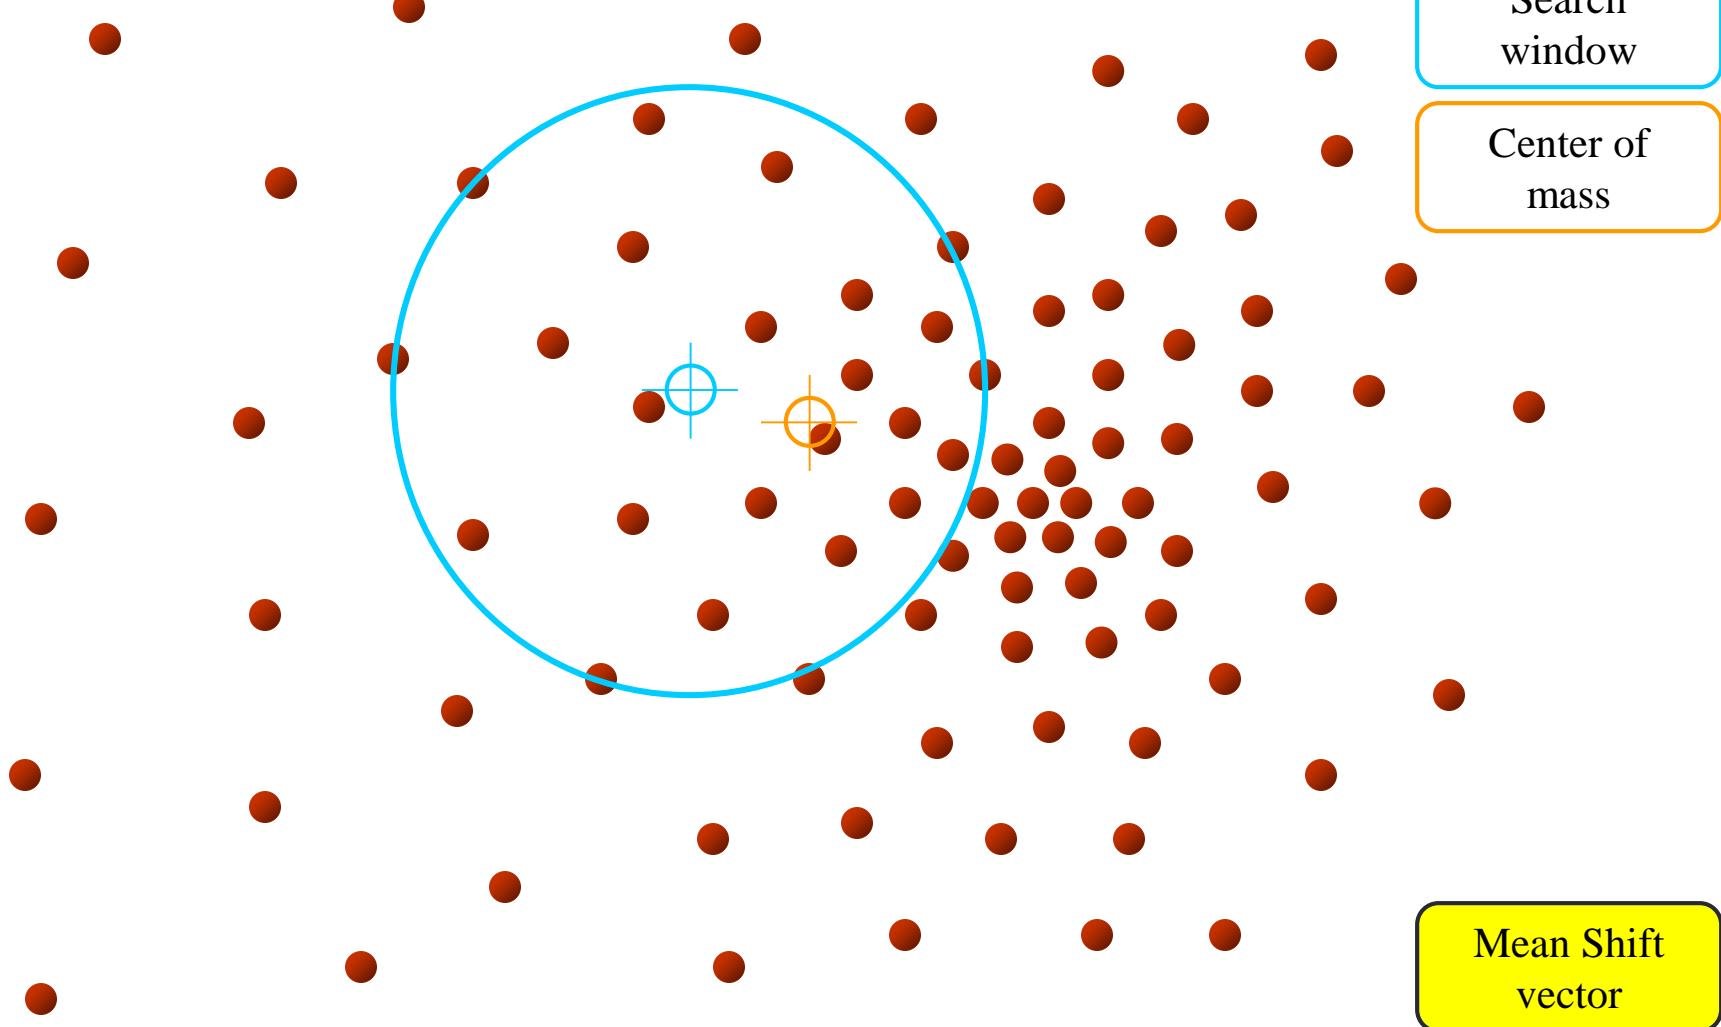
Feature space
($L^*u^*v^*$ color values)



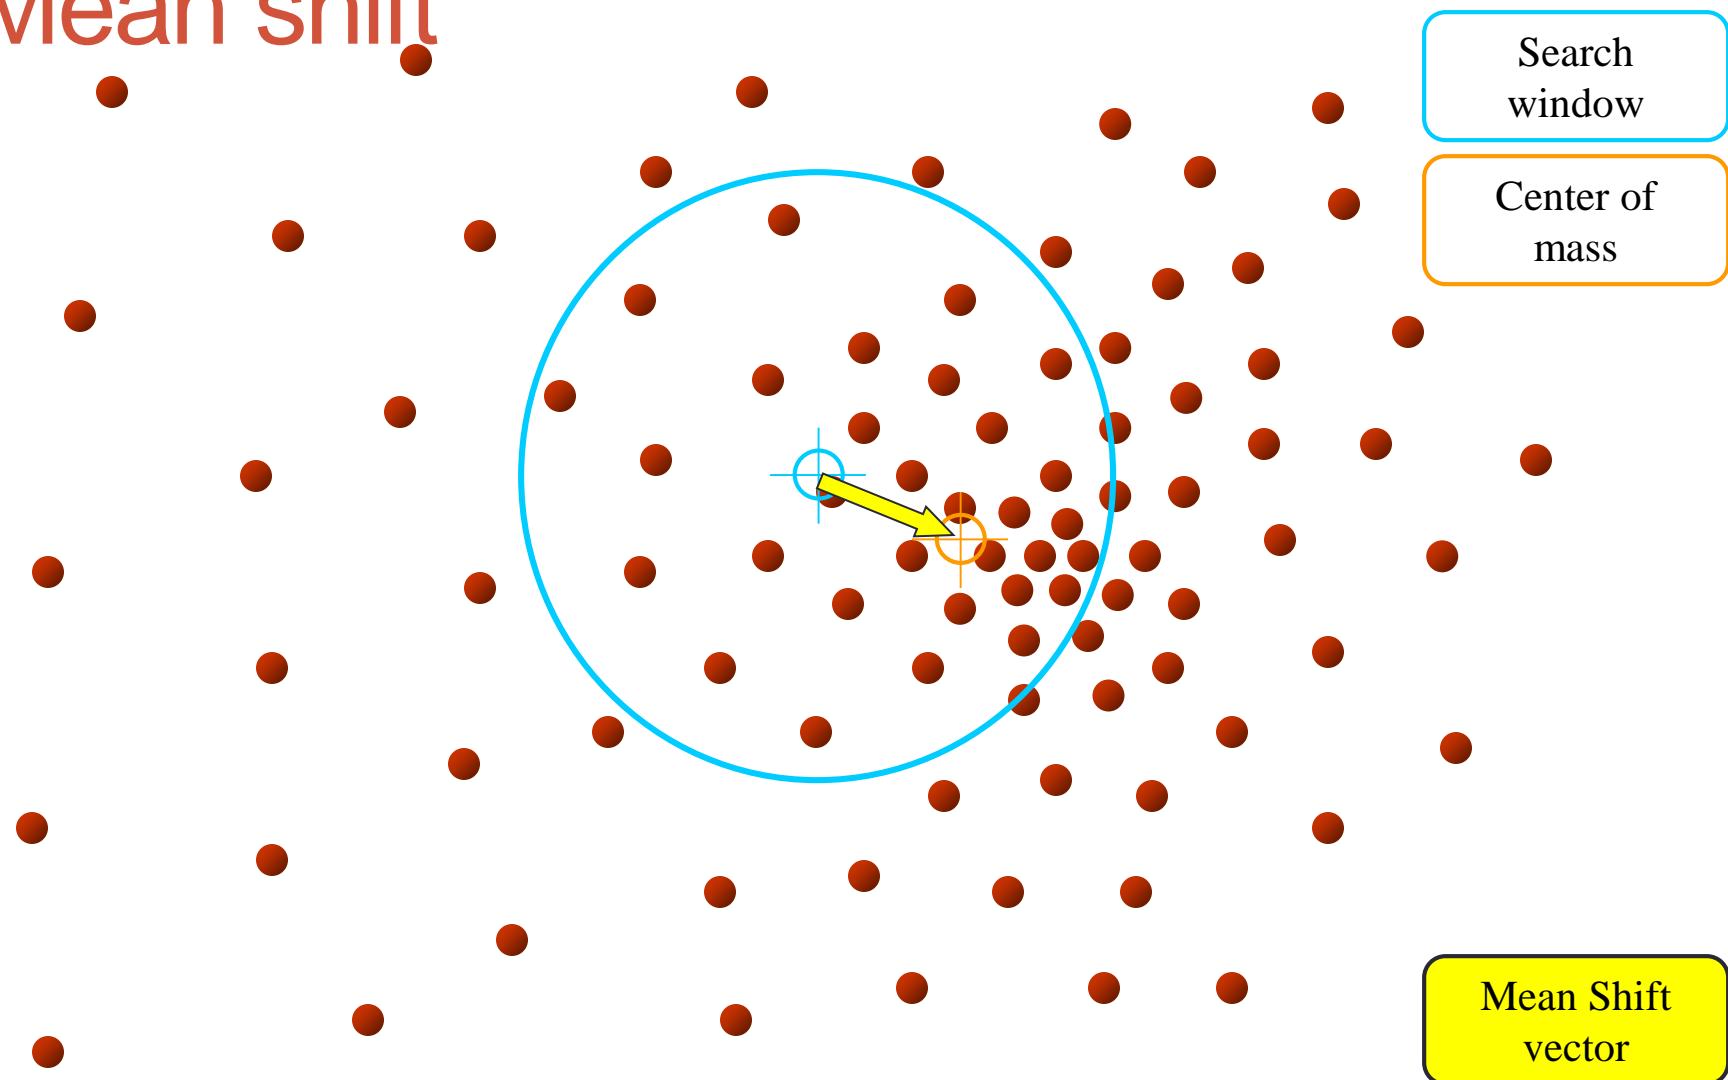
Mean shift



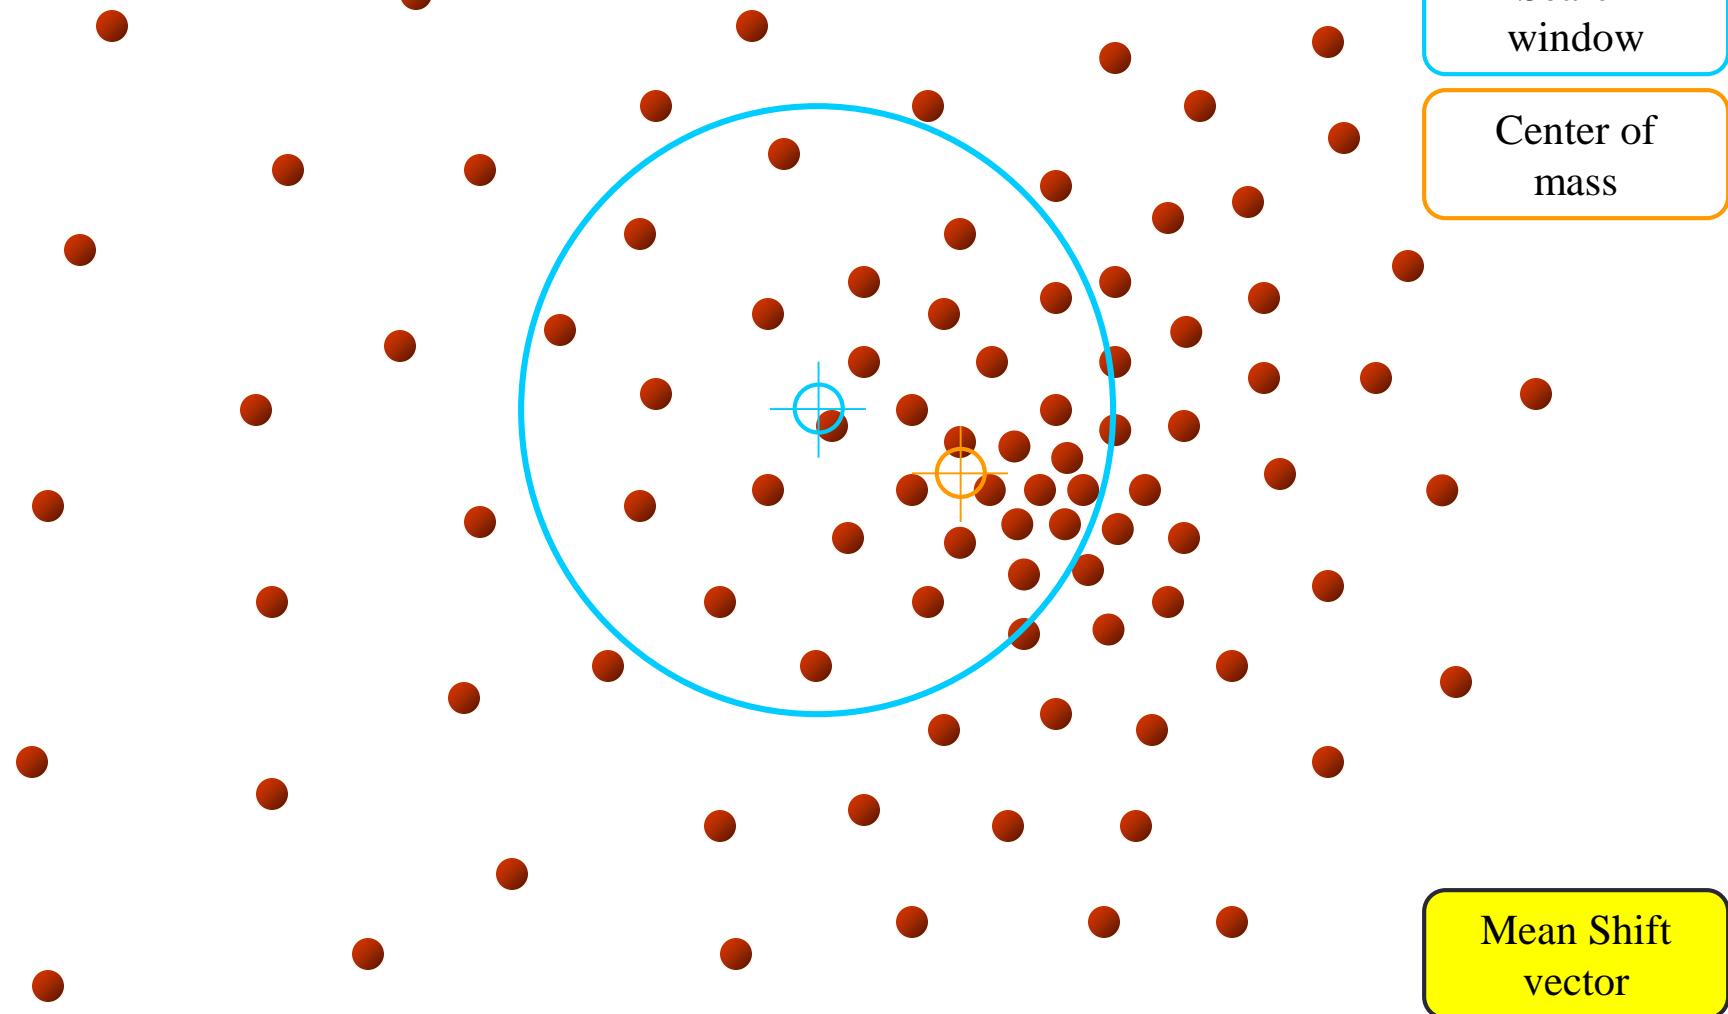
Mean shift



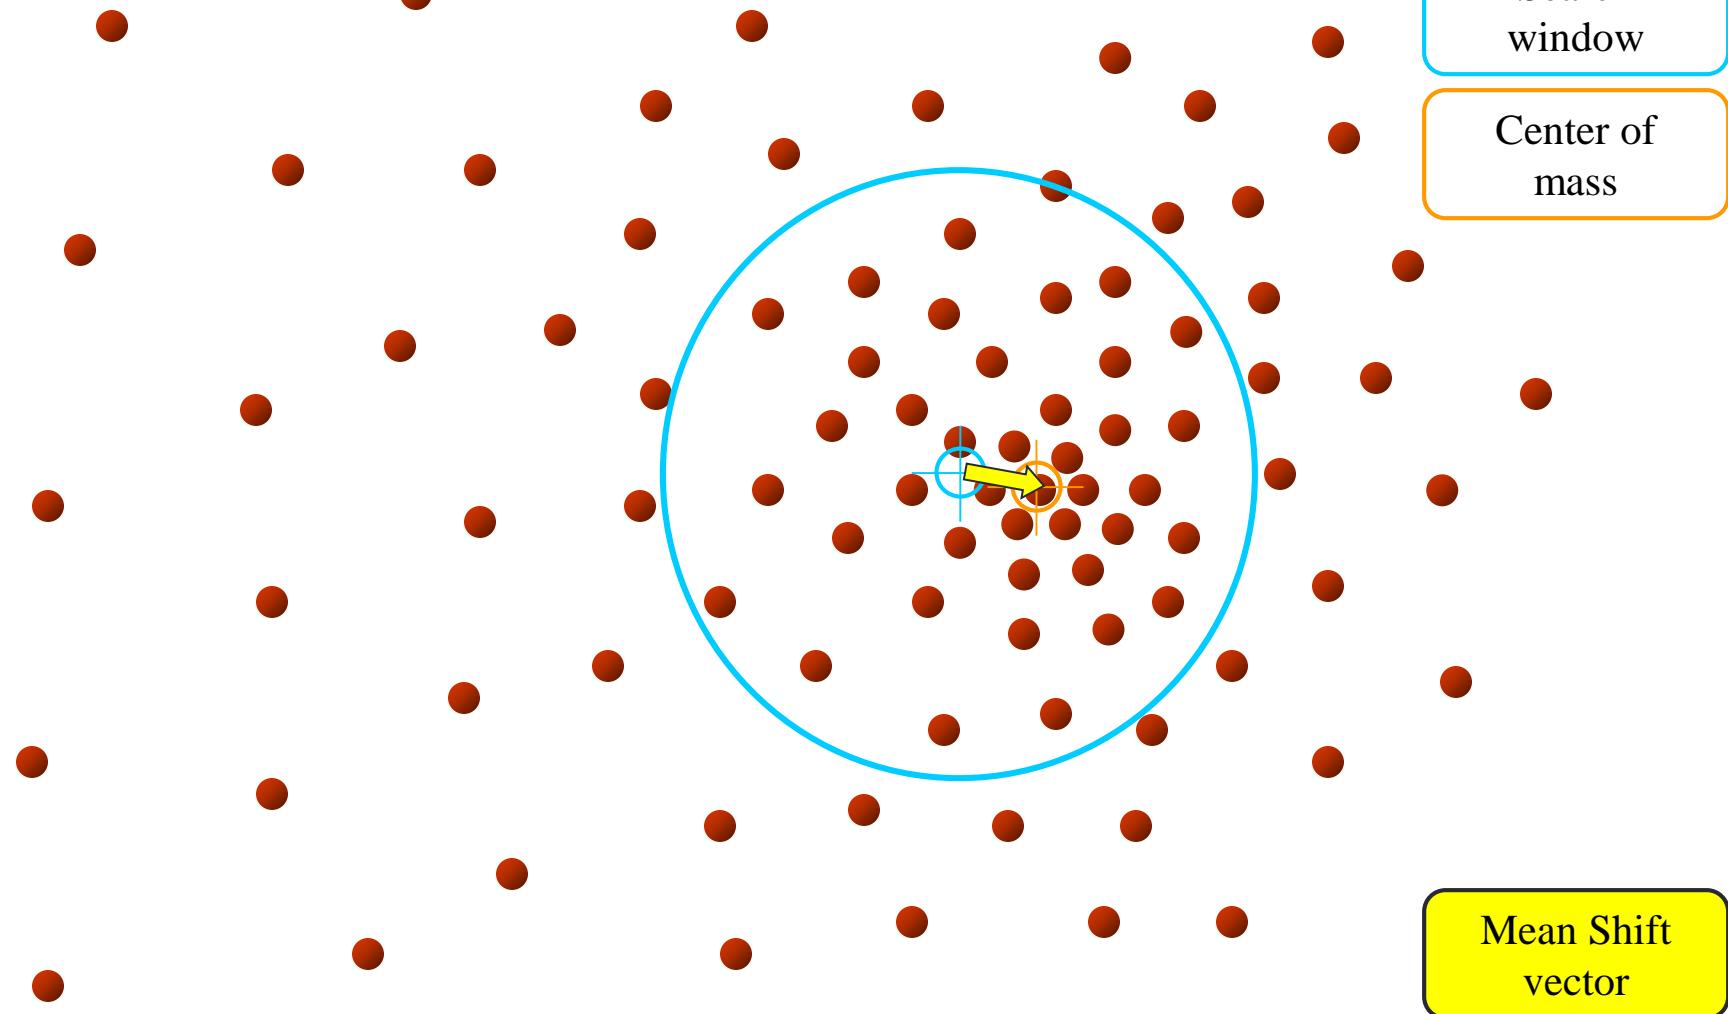
Mean shift



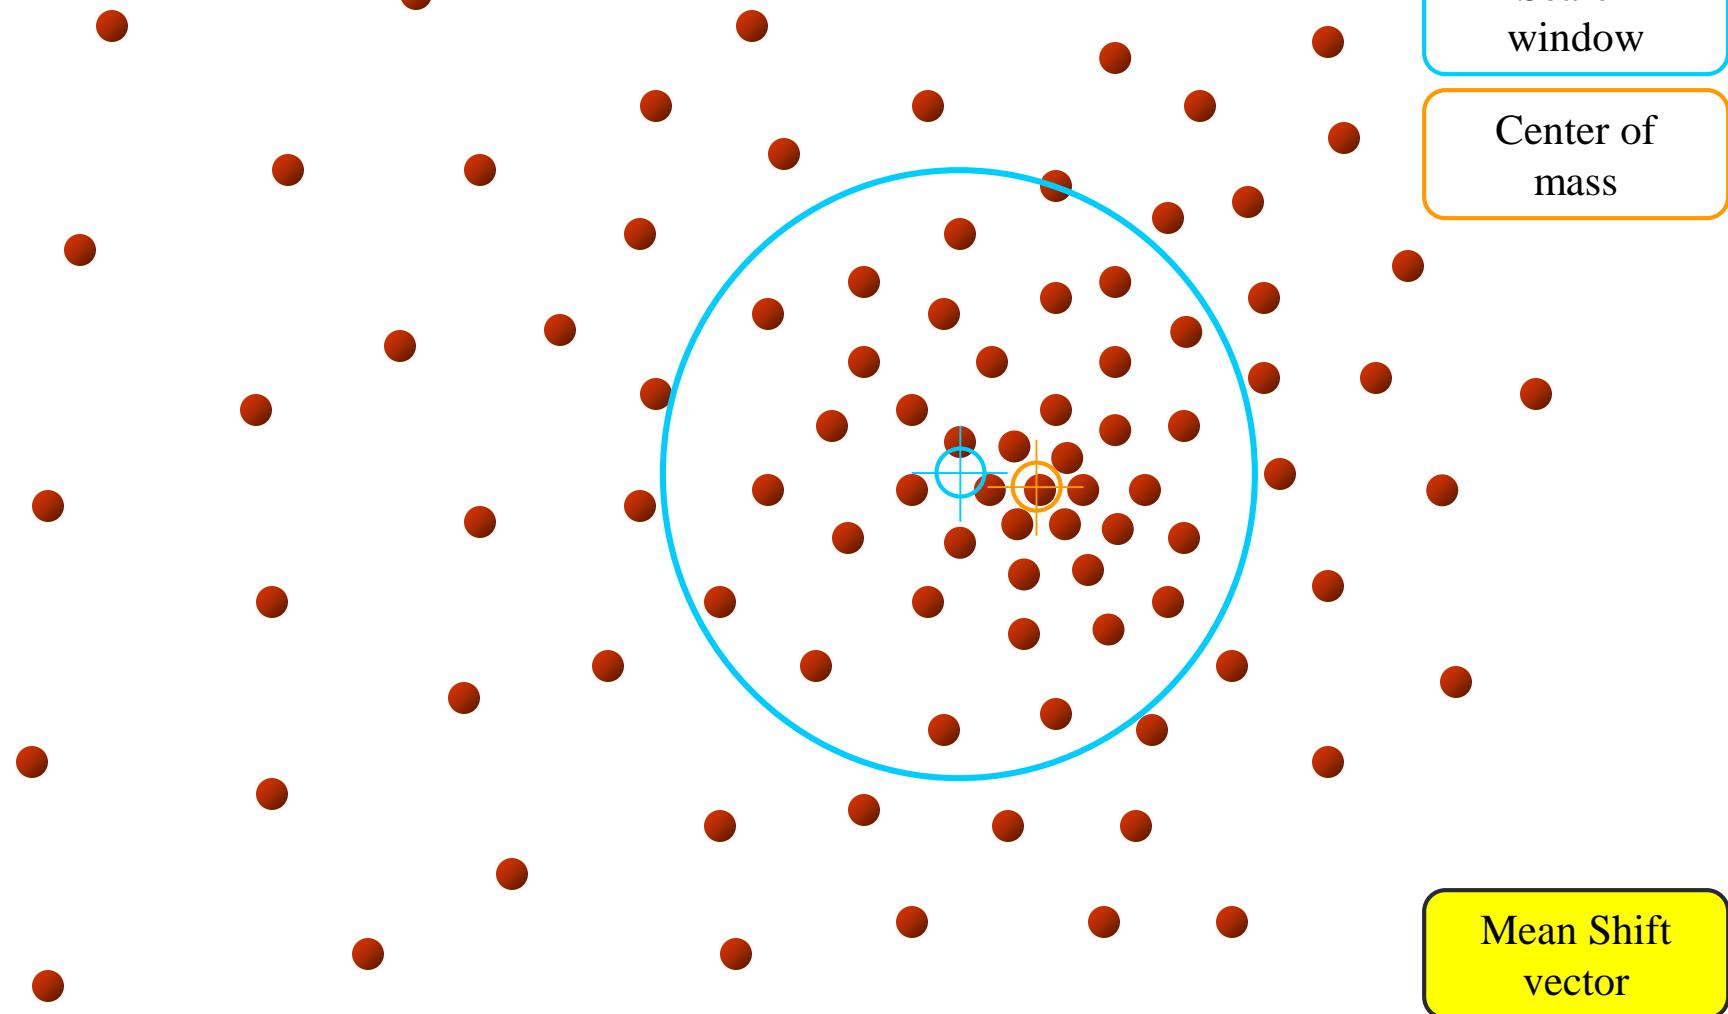
Mean shift



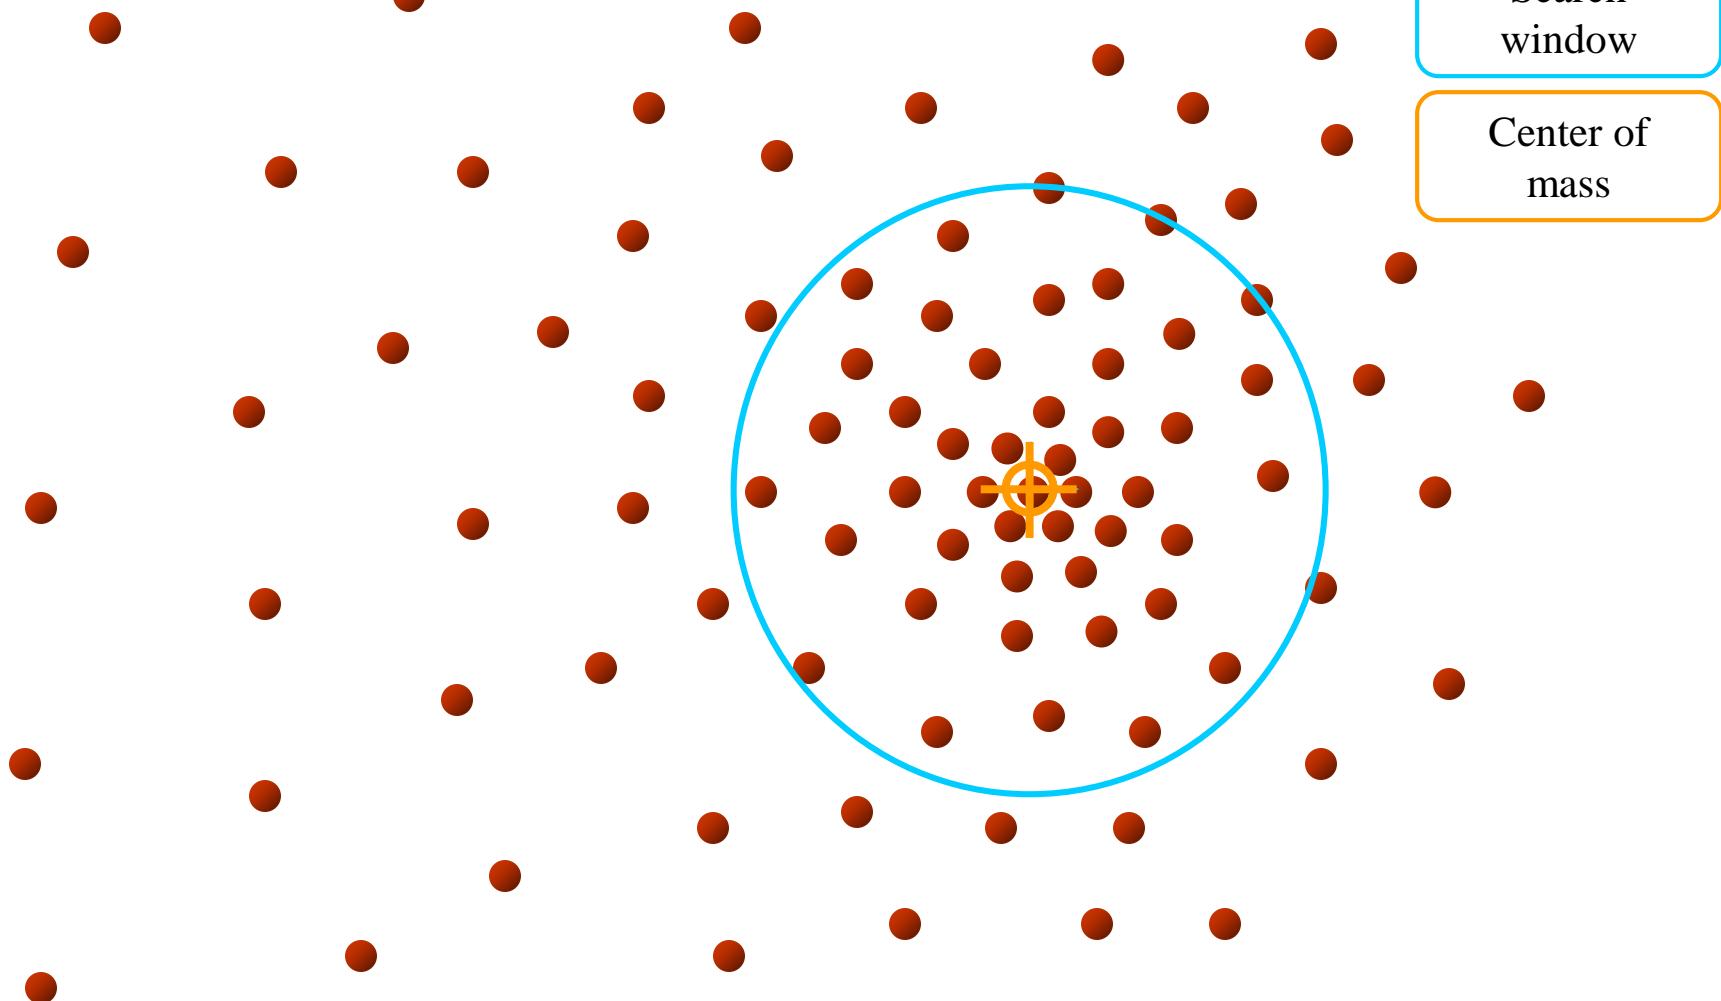
Mean shift



Mean shift



Mean shift

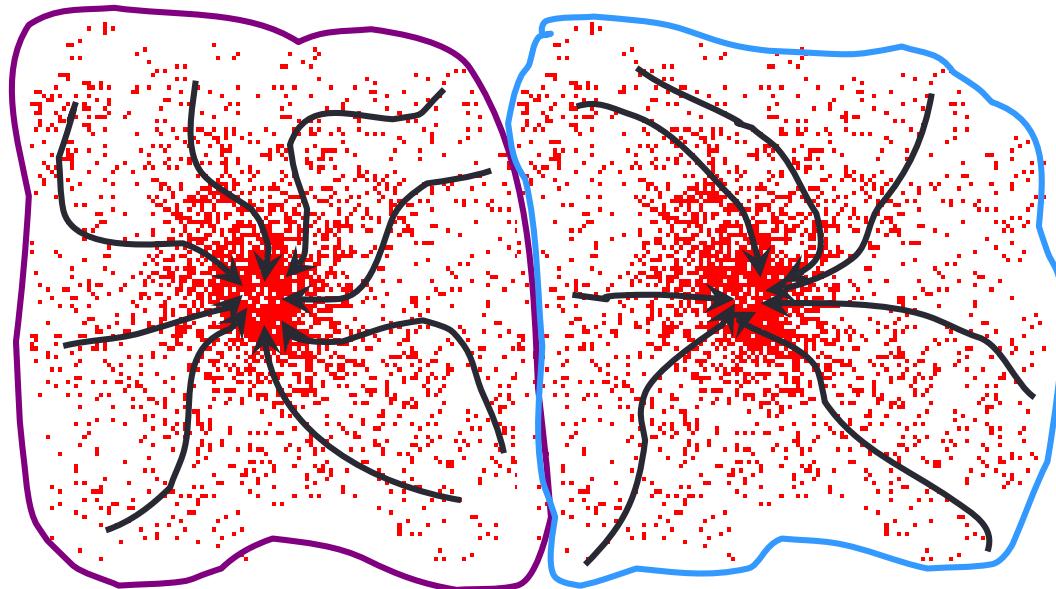


Search
window

Center of
mass

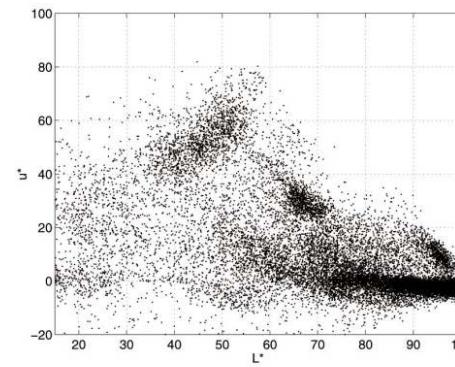
Mean shift clustering

- Cluster: all data points in the attraction basin of a mode
- Attraction basin: the region for which all trajectories lead to the same mode

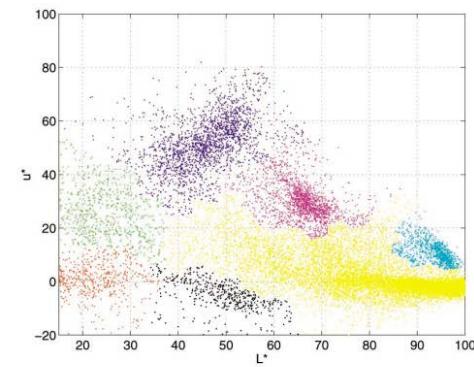


Mean shift clustering/segmentation

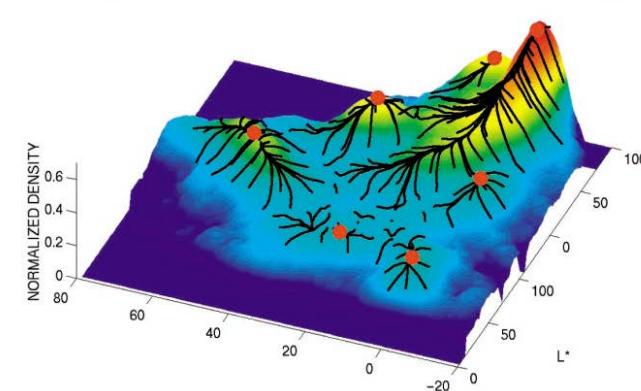
- Find features (color, gradients, texture, etc)
- Initialize windows at individual feature points (pixels)
- Perform mean shift for each window (pixel) until convergence
- Merge windows (pixels) that end up near the same “peak” or mode



(a)



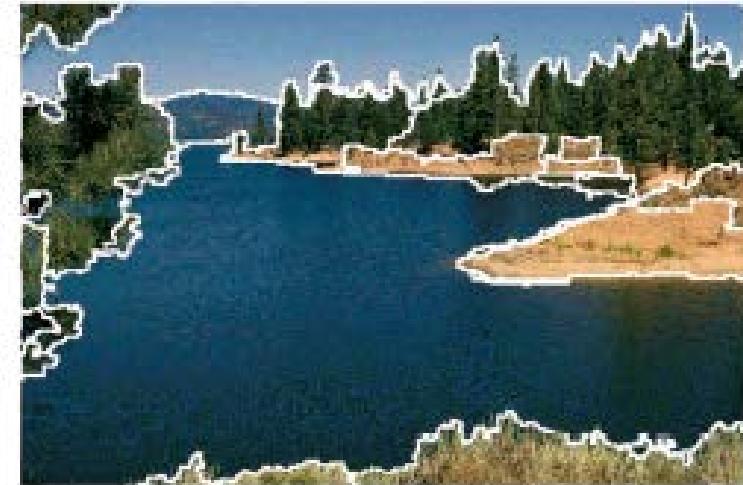
(b)



Mean shift segmentation results



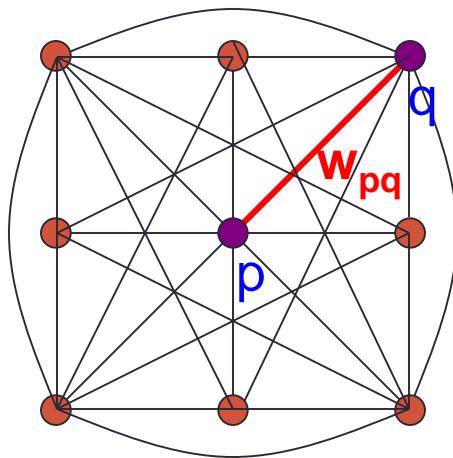
Mean shift segmentation results



Mean shift

- Pros:
 - Does not assume shape on clusters
 - One parameter choice (window size)
 - Generic technique
 - Find multiple modes
- Cons:
 - Selection of window size
 - Does not scale well with dimension of feature space

Images as graphs



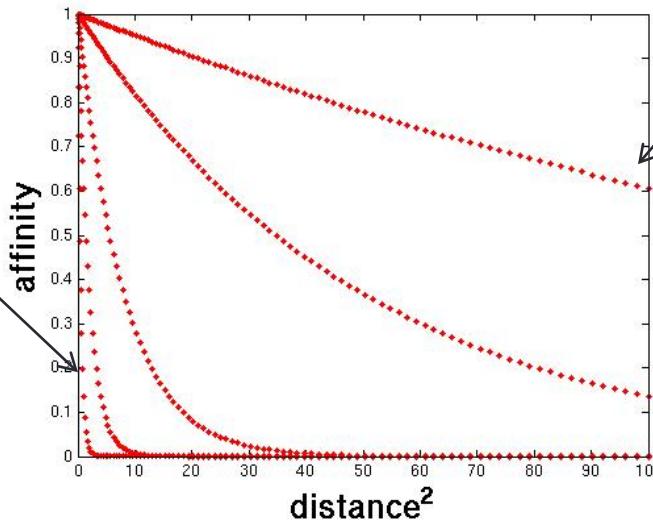
- *Fully-connected* graph
 - node (vertex) for every pixel
 - link between *every* pair of pixels, \mathbf{p}, \mathbf{q}
 - affinity weight $w_{\mathbf{pq}}$ for each link (edge)
 - $w_{\mathbf{pq}}$ measures *similarity*
 - similarity is *inversely proportional* to difference (in color and position...)

Measuring affinity

- One possibility:

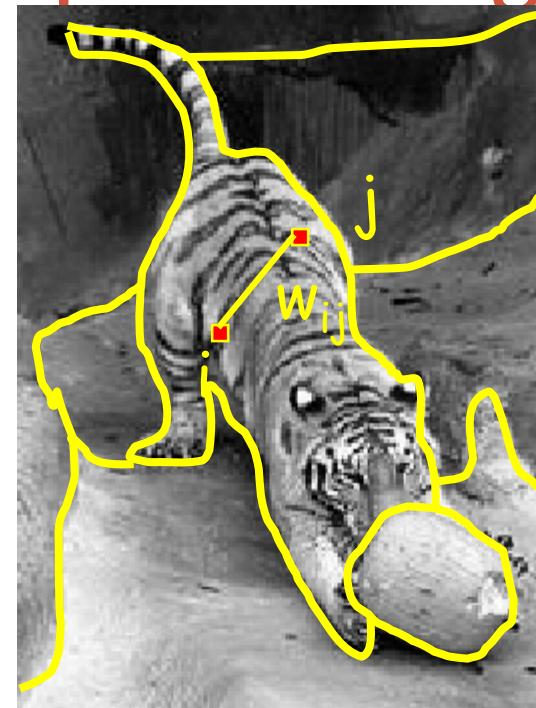
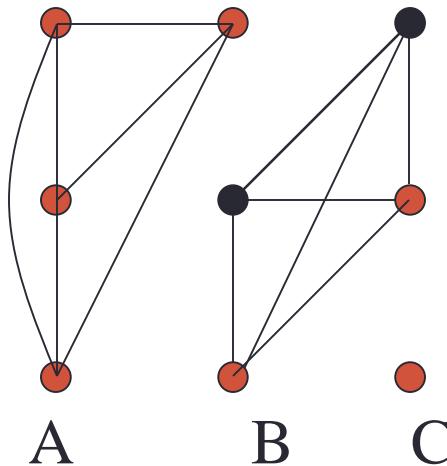
$$aff(\mathbf{x}_i, \mathbf{x}_j) = \exp\left(-\frac{1}{2\sigma^2} \text{dist}(\mathbf{x}_i, \mathbf{x}_j)^2\right)$$

Small sigma:
group only
nearby points



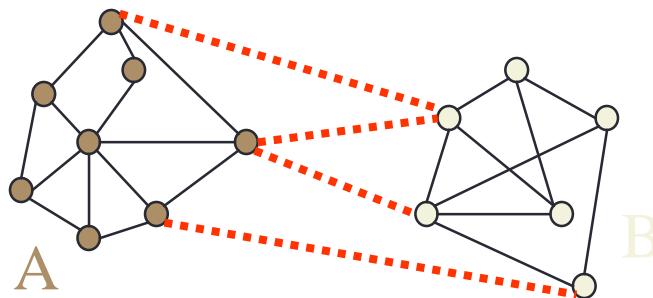
Large sigma:
group distant
points

Segmentation by graph partitioning



- Break Graph into Segments
 - Delete links that cross between segments
 - Easiest to break links that have low affinity
 - similar pixels should be in the same segments
 - dissimilar pixels should be in different segments

Graph cut

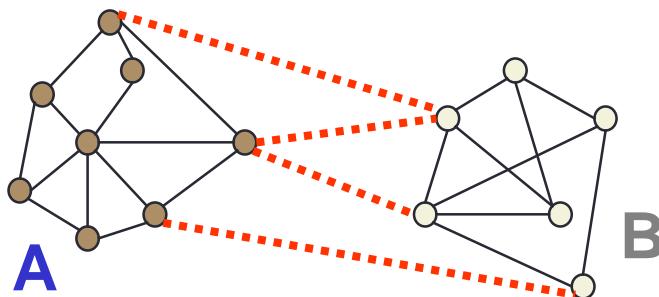


- Set of edges whose removal makes a graph disconnected
- Cost of a cut: sum of weights of cut edges

$$\text{cut}(A, B) = \sum_{p \in A, q \in B} w_{p,q}$$

- A graph cut gives us a segmentation
 - What is a “good” graph cut and how do we find one?

Cuts in a graph: Min cut



$$\text{cut}(A, B) = \sum_{p \in A, q \in B} w_{p,q}$$

Find minimum cut

- gives you a segmentation
- fast algorithms exist for doing this (we may see this...)

Minimum cut

- Problem with minimum cut:

Weight of cut proportional to number of edges in the cut;
tends to produce small, isolated components.

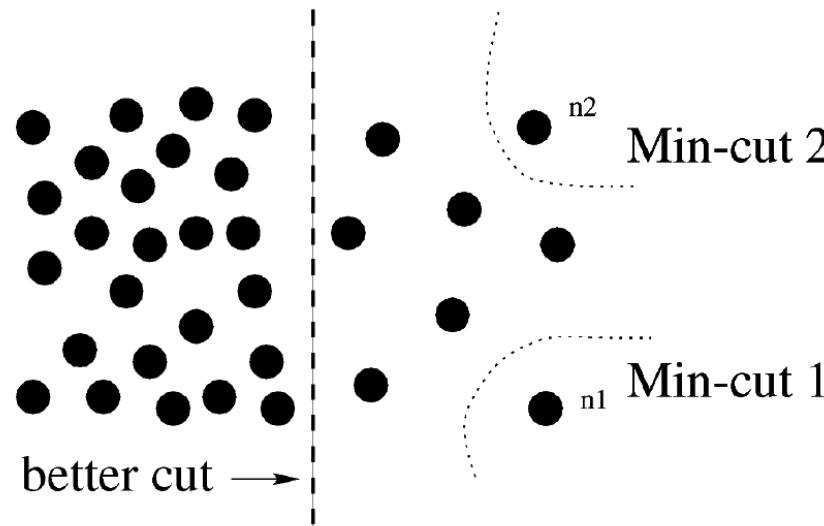
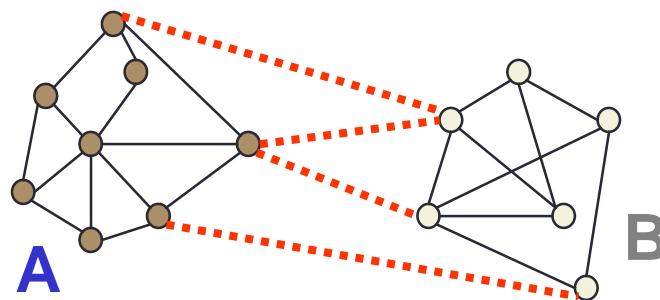


Fig. 1. A case where minimum cut gives a bad partition.

Cuts in a graph: Normalized cut



Normalized Cut

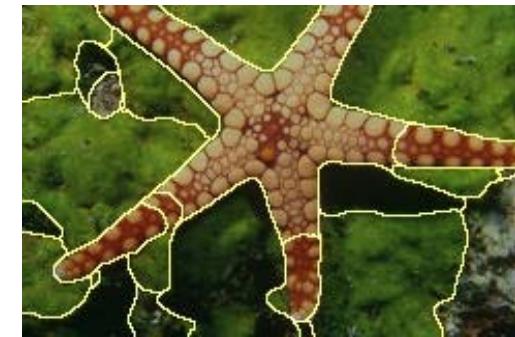
- fix bias of Min Cut by **normalizing** for size of segments:

$$Ncut(A, B) = \frac{cut(A, B)}{assoc(A, V)} + \frac{cut(A, B)}{assoc(B, V)}$$

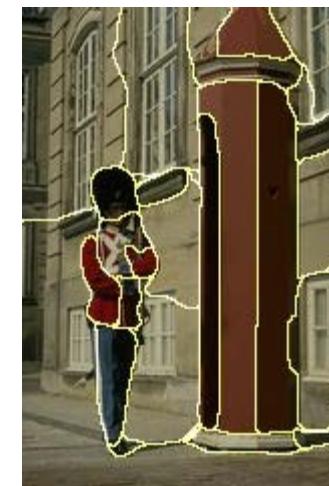
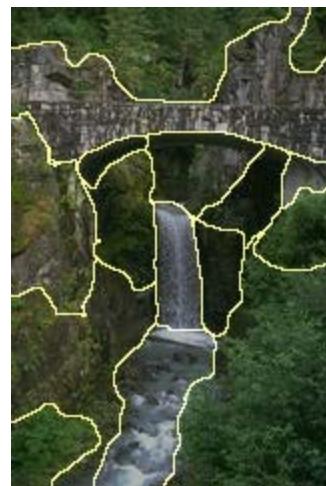
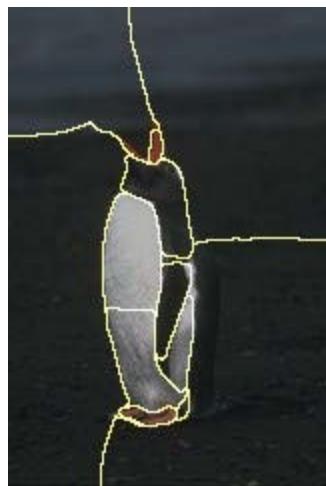
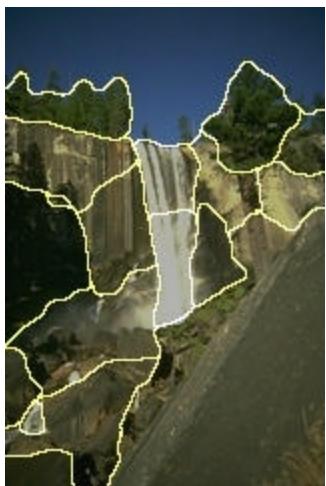
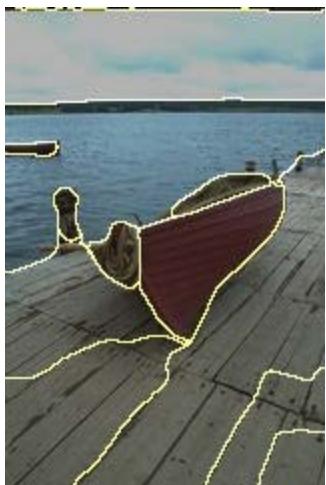
$assoc(A, V)$ = sum of weights of all edges that touch A

- Ncut value small when we get two clusters with many edges with high weights, and few edges of low weight between them
- Approximate solution for minimizing the Ncut value : generalized eigenvalue problem.

Example results



Results: Berkeley Segmentation Engine



<http://www.cs.berkeley.edu/~fowlkes/BSE/>

Normalized cuts: pros and cons

Pros:

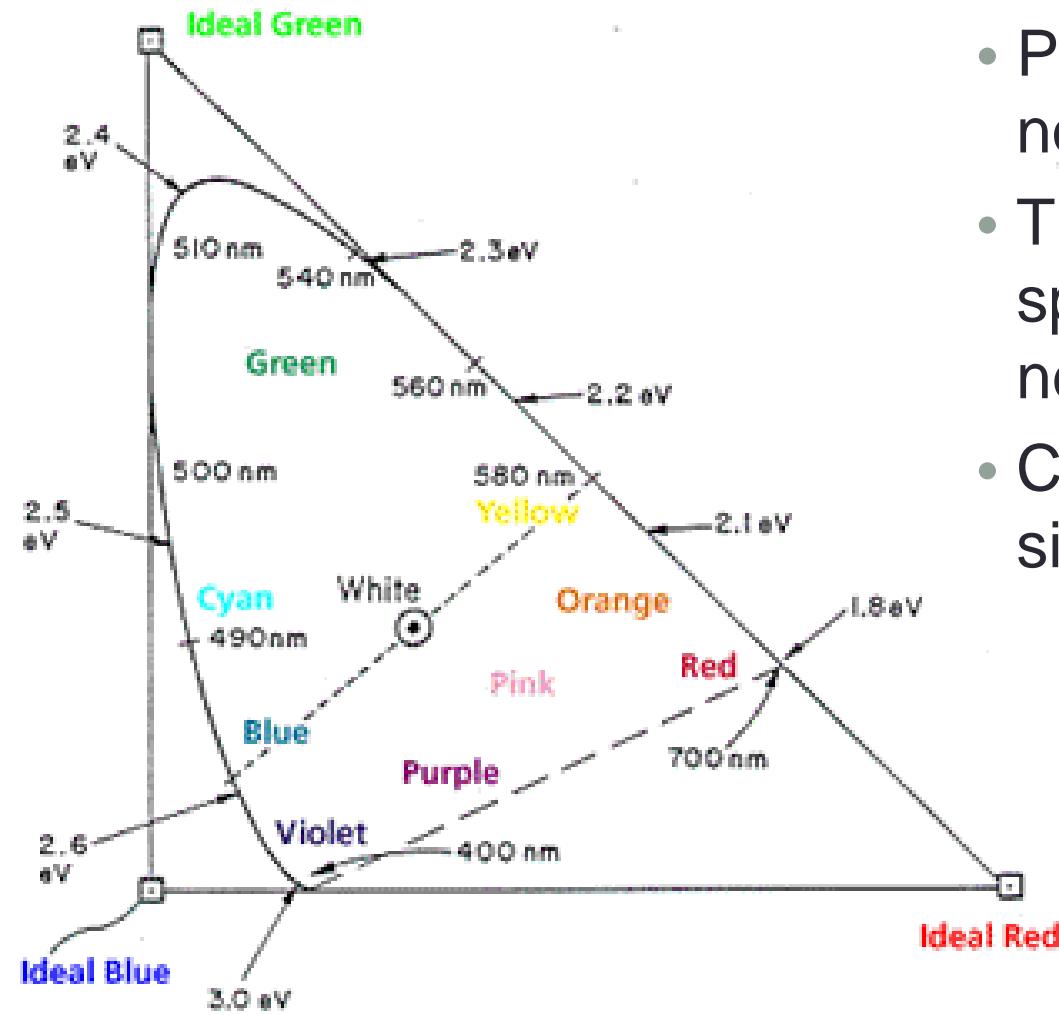
- Generic framework, flexible to choice of function that computes weights (“affinities”) between nodes
- Does not require model of the data distribution

Cons:

- Time complexity can be high
 - Dense, highly connected graphs → many affinity computations
 - Solving eigenvalue problem
- Preference for balanced partitions

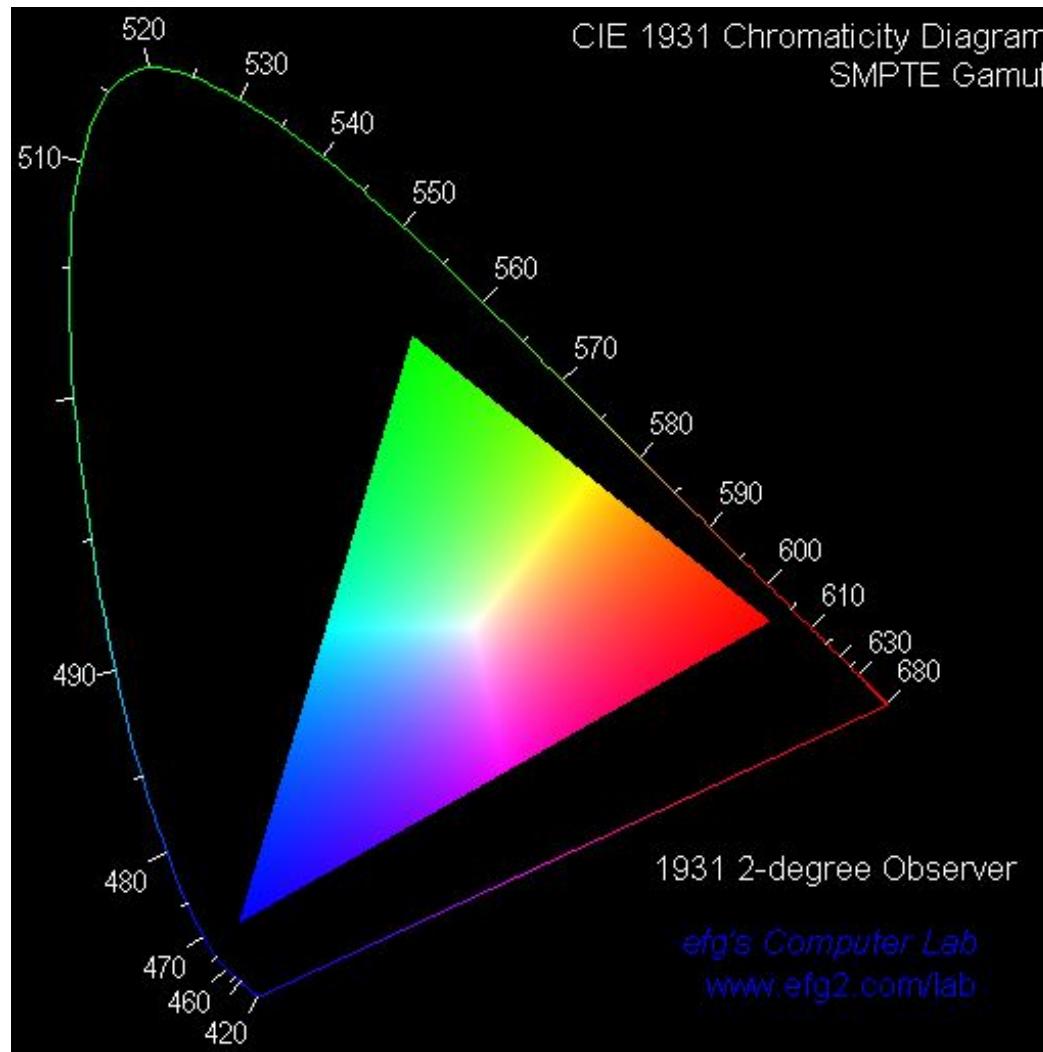
The end...

Geometry of Color (CIE)



- Perceptual color spaces are non-convex
- Three primaries can span the space, but weights may be negative.
- Curved outer edge consists of single wavelength primaries

RGB Color Space



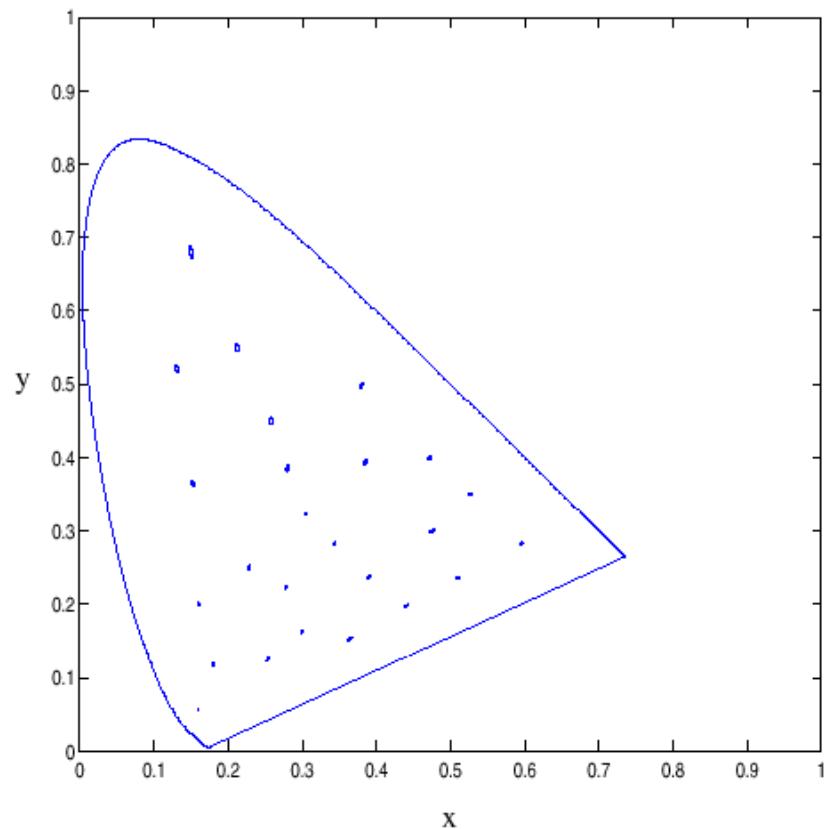
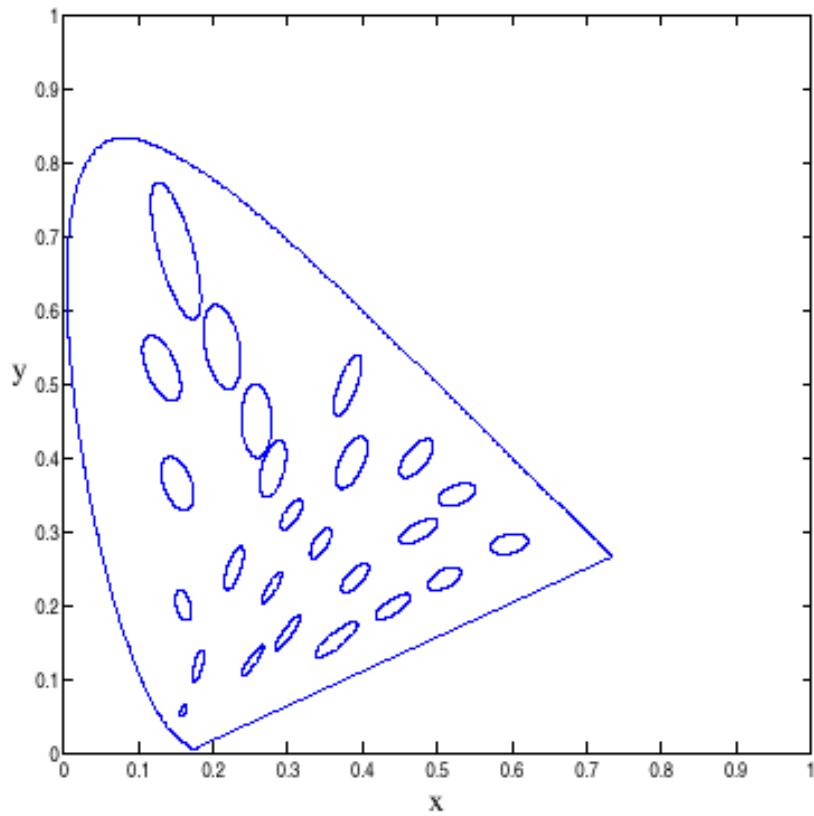
Many colors cannot be represented
(phosphor limitations)

Uniform color spaces

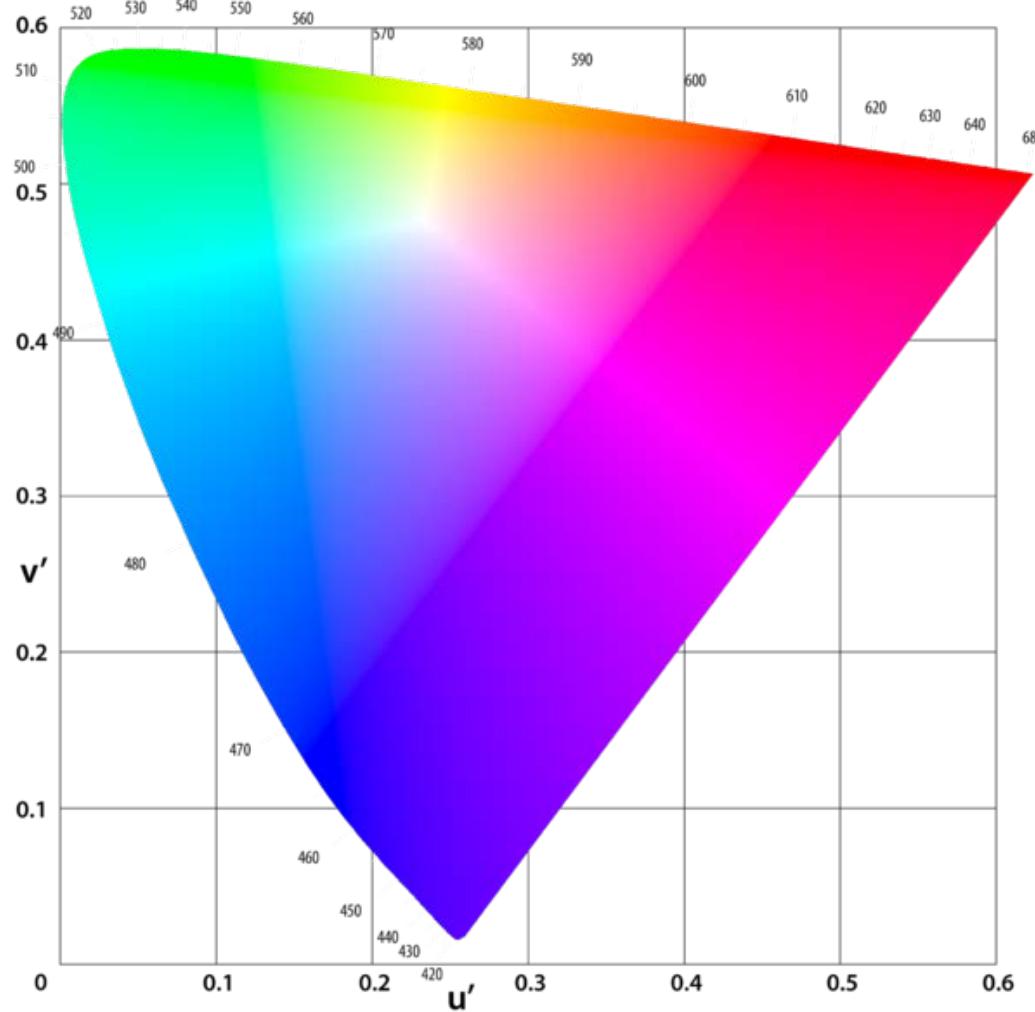
- McAdam ellipses (next slide) demonstrate that differences in x,y are a poor guide to differences in color
- Construct color spaces so that differences in coordinates are a good guide to differences in color.

McAdam ellipses

Figures courtesy of
D. Forsyth

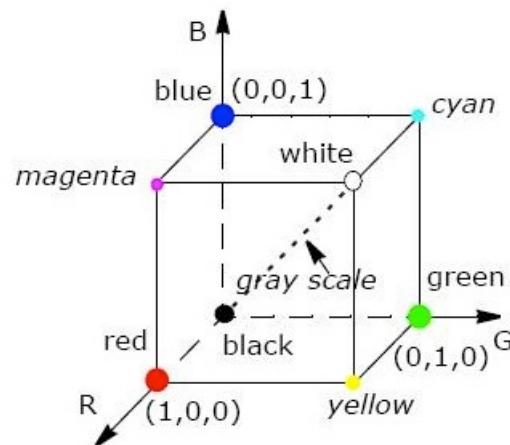


LUV Color Space

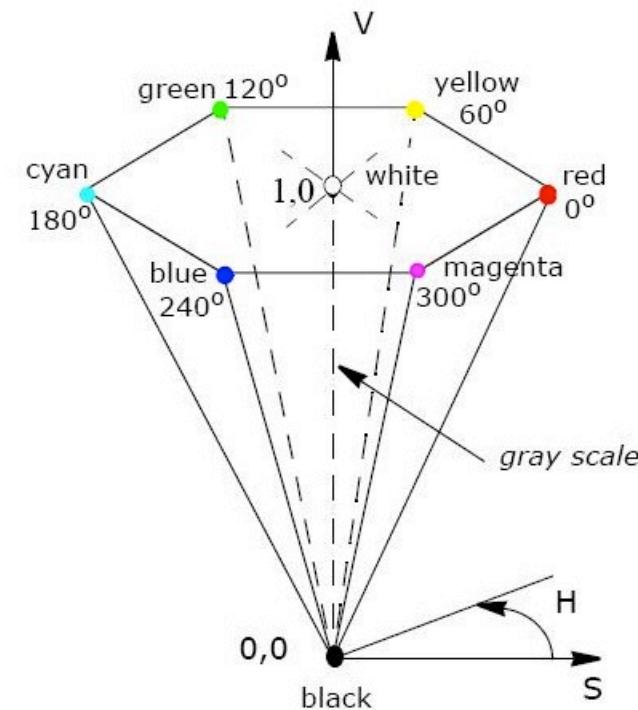


HSV Color Space

- RGB

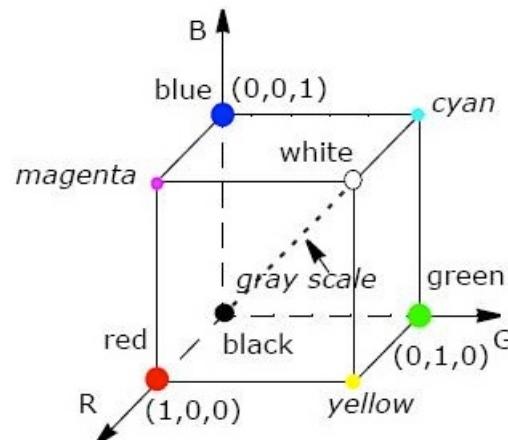


- HSV

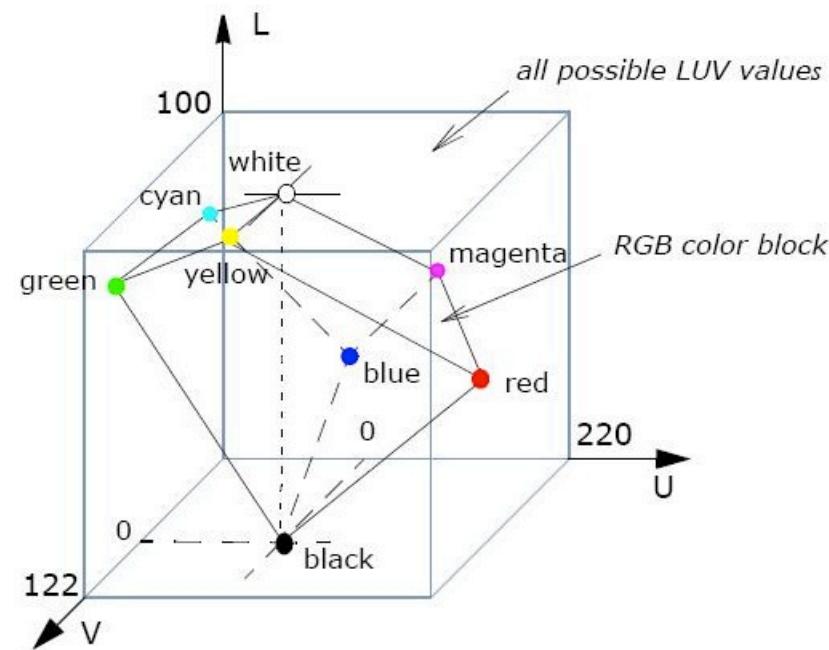


LUV Color Space

- RGB



- LUV



More info see: http://software.intel.com/sites/products/documentation/hpc/ipp/ippi_ch6/ch6_color_models.html